Flare32 CPU

FL4SHK

April 9, 2023

Table of Contents

Introduction

2.1 Registers

There are sixteen general-purpose registers: `r0`, `r1`, `r2`, ..., `r11`, `r12`, `lr`, `fp`, `sp`. Each register is 32 bits long. For special purpose registers, there are also `pc`, the program counter (which is 32 bits long), and the `flags`. Also there are the interrupts-related registers: `ids` (the destination to go to upon an interrupt happening), `ira` (the program counter value to return to after an interrupt) and `ie` (whether or not interrupts are enabled). Two more registers are `hi` and `lo`, which are used as the high 32 bits and low 32 bits of the result of a 32 by 32 -> 64 multiplication, or as the high 32 bits and low 32 bits of the result of a 64 by 64 -> 64 division. Here are the flags:

Table 1: The Flags

Zero (Z)	Carry (C)	oVerflow (V)	Negative (N)

Instruction Set

3.1 Instruction Group 0: `pre` and `lpre`

For `pre`, the following encoding is used, with each character representing one bit:

`0000 iiii iiii iiii`, where

- `i` is a 12-bit constant.

For `lpre`, the following encoding is used, with each character representing one bit:
`0001 0iii iiii iiii iiii iiii iiii iiii`, where

- `i` is a 27-bit constant.

`pre` and `lpre` are mechanisms by which immediates larger than normal can be used, essentially acting like variable width instructions.

There is no mechanism in the assembly language itself to use `pre` or `lpre` as instructions. Instead, it is expected that the assembler or linker will be the one to insert `pre` or `lpre` as needed if an immediate is too large for a particular instruction.

For non-branch instructions:

- If `pre` is used, the immediate field of the `pre` instruction will form bits `[16:5]` of the immediate of the next non-`index` instruction. The 17-bit immediate will then be sign-extended to 32 bits.

- On the other hand, if `lpre` is used, the immediate field of the `lpre` instruction will form bits `[31:5]` of the immediate of the next non-`index` instruction.

For branch instructions (group 3):

- if `pre` is used, the immediate field of the `pre` instruction will form bits `[23:9]` of the immediate of the next non-`index` instruction. The 23-bit immediate will then be sign-extended to 32 bits.

- On the other hand, if `lpre` is used, the immediate field of the `lpre` instruction will form bits `[31:9]` of the immediate of the next non-`index` instruction.

Instructions that use an immediate value, other than `pre`, `lpre`, and relative branches normally have 5-bit immediates. Relative branches normally have 9-bit branch offsets. if `pre` or `lpre` is used with a relative branch, bits `[9:5]` of the branch offset that is encoded directly into a relative branch instruction will be ignored, and the immediate field of the `pre` or `lpre` instruction will be used to determine bits `[9:5]` of the branch offset..

3.1.1 Handling of `pre`, `lpre`, and `index`

When a `pre` or `lpre` instruction is found, `pre` or `lpre` will be considered to be "in effect". This condition lasts for one or two instructions after the `pre` or `lpre` instruction, depending on whether or not `index` was in effect.

`index` is an instruction (defined later) that allows a subsequent load or store instruction to perform `base_reg + index_reg` indexing. When an `index` instruction is found, it is considered to be in effect. Like `pre` and `lpre`, `index` is inserted automatically by the assembler.

`index` can be combined with `pre` or `lpre`, and it does not matter whether `index` or the `pre`/`lpre` instruction came first.

When `pre`, `lpre`, or `index` is in effect, interrupts will not be serviced.

Pseudo code for handling the how `pre`, `lpre`, and `index` are to be handled regarding whether or not they are "in effect" is as follows:

```
function handle_pre_lpre_index(input in, output out) {
    if (in.instruction.is_pre()) {
        if (in.state.pre.have || in.state.lpre.have) {
            out.state.can_service_interrupts = false;

            // Instruction was a NOP
            out.state.pre.have = false;
            out.state.lpre.have = false;
            out.state.index.have = false;
            out.instruction = nop;
        } else {
            if (in.state.index.have) {
                out.state.can_service_interrupts = false;
            } else {
                out.state.can_service_interrupts = true;
            }
            out.state.pre.set_have(true);
            out.instruction = in.instruction;
        }
    } else if (in.instruction.is_lpre()) {
        if (in.state.pre.have || in.state.lpre.have) {
            out.state.can_service_interrupts = false;

            // Instruction was a NOP
            out.state.pre.have = false;
            out.state.lpre.have = false;
            out.state.index.have = false;
            out.instruction = nop;
        } else {
            if (in.state.index.have) {
                out.state.can_service_interrupts = false;
```

```
            } else {
                out.state.can_service_interrupts = true;
            }
            out.state.lpre.set_have(true);
            out.instruction = in.instruction;
        }
    } else if (in.instruction.is_index()) {
        if (in.state.index.have) {
            out.state.can_service_interrupts = false;

            // Instruction was a NOP
            out.state.pre.have = false;
            out.state.lpre.have = false;
            out.state.index.have = false;
            out.instruction = nop;
        } else {
            if (in.state.pre.have || in.state.lpre.have) {
                out.state.can_service_interrupts = false;
            } else {
                out.state.can_service_interrupts = true;
            }
            out.state.index.have = true;
            out.instruction = in.instruction;
        }
    } else {
        if (
            in.state.pre.have || in.state.lpre.have || in.state.index.have
        ) {
            out.state.can_service_interrupts = false;
        } else {
            out.state.can_service_interrupts = true;
        }
        out.instruction = instruction;

        // Whenever we see an instruction other than pre, lpre, or index,
        // that means that those instructions stop being "in effect".
        out.state.pre.have = false;
        out.state.lpre.have = false;
        out.state.index.have = false;
    }
}
```

3.2   Instruction Group 1

The following encoding is used, with each character representing one bit:
`001i iiii oooo aaaa`, where

- `i` is a 5-bit sign-extended immediate, and is denoted `simm`

- `a` encodes register `rA`

- `o` is the opcode

Here is a list of instructions from this encoding group.

- Opcode 0x0: **`add rA, #simm`**

- Opcode 0x1: **`add rA, pc, #simm`**

- Opcode 0x2: **`add rA, sp, #simm`**

- Opcode 0x3: **`add rA, fp, #simm`**

- Opcode 0x4: **`cmp rA, #simm`**

  - Note: Compare `rA` to `simm`.
  - Affectable flags: `Z, C, V, N`

- Opcode 0x5: **`cpy rA, #simm`**

  - Note: Copy an immediate value into `rA`

- Opcode 0x6: **`lsl rA, #simm`**

  - Note: Logical shift left

- Opcode 0x7: **`lsr rA, #simm`**

  - Note: Logical shift right

- Opcode 0x8: **`asr rA, #simm`**

  - Note: Arithmetic shift right

- Opcode 0x9: **`and rA, #simm`**

  - Note: Bitwise AND

- Opcode 0xa: **`orr rA, #simm`**

  - Note: Bitwise OR

- Opcode 0xb: **`xor rA, #simm`**

  - Note: Bitwise XOR

- Opcode 0xc: **`zeb rA`**

  - Effect: Set `rA[31:8]` to zero.

- Opcode `0xd`: **`zeh rA`**

  - Effect: Set `rA[31:16]` to zero.

- Opcode `0xe`: **`seb rA`**

  - Effect: Sign-extend `rA[7:0]` to 32 bits, then copy that value to `rA`

- Opcode `0xf`: **`seh rA`**

  - Effect: Sign-extend `rA[15:0]` to 32 bits, then copy that value to `rA`

3.3   Instruction Group 2

The following encoding is used, with each character representing one bit:

`010f oooo cccc aaaa`, where

- `o` is the opcode

- `c` encodes register `rC`

- `a` encodes register `rA`

- `f` is encoded as `0` if this instruction cannot affect flags and encoded `1` if this instruction is permitted to affect flags. Note that `cmp` is permitted to affect flags regardless of this bit.

Here is a list of instructions from this encoding group.

- Opcode `0x0`: **`add rA, rC`**

  - Mnemonic for when flags not affected: **`add`**
  - Mnemonic for when flags affected: **`add.f`**
  - Affectable flags: `Z`, `C`, `V`, `N`

- Opcode `0x1`: **`sub rA, rC`**

  - Mnemonic for when flags not affected: **`sub`**
  - Mnemonic for when flags affected: **`sub.f`**
  - Affectable flags: `Z`, `C`, `V`, `N`

- Opcode `0x2`: **`add rA, sp, rC`**

  - Mnemonic for when flags not affected: **`add`**
  - Mnemonic for when flags affected: **`add.f`**
  - Affectable flags: `Z`, `C`, `V`, `N`

7

- Opcode `0x3` `add rA, fp, rC`

  - Mnemonic for when flags not affected: `add`
  - Mnemonic for when flags affected: `add.f`
  - Affectable flags: `Z, C, V, N`

- Opcode `0x4`: `cmp rA, rC`

  - Note: Compare `rA` to `rC`. `cmp` is *always* able to affect flags, independent of the encoded `f` bit of the instruction.
  - Affectable flags: `Z, C, V, N`

- Opcode `0x5`: `cpy rA, rC`

  - Mnemonic for when flags not affected: `cpy`
  - Mnemonic for when flags affected: `cpy.f`
  - Note: Copy `rC` into `rA`
  - Affectable flags: `Z, N`

- Opcode `0x6`: `lsl rA, rC`

  - Mnemonic for when flags not affected: `lsl`
  - Mnemonic for when flags affected: `lsl.f`
  - Note: Logical shift left
  - Affectable flags: `Z, N`

- Opcode `0x7`: `lsr rA, rC`

  - Mnemonic for when flags not affected: `lsr`
  - Mnemonic for when flags affected: `lsr.f`
  - Note: Logical shift right
  - Affectable flags: `Z, N`

- Opcode `0x8`: `asr rA, rC`

  - Mnemonic for when flags not affected: `asr`
  - Mnemonic for when flags affected: `asr.f`
  - Note: Arithmetic shift right
  - Affectable flags: `Z, N`

- Opcode `0x9`: `and rA, rC`

  - Mnemonic for when flags not affected: `and`
  - Mnemonic for when flags affected: `and.f`
  - Note: Bitwise AND

– Affectable flags: `Z`, `N`

- Opcode `0xa`: **`orr rA, rC`**

  – Mnemonic for when flags not affected: `orr`

  – Mnemonic for when flags affected: `orr.f`

  – Note: Bitwise OR

  – Affectable flags: `Z`, `N`

- Opcode `0xb`: **`xor rA, rC`**

  – Mnemonic for when flags not affected: `xor`

  – Mnemonic for when flags affected: `xor.f`

  – Note: Bitwise XOR

  – Affectable flags: `Z`, `N`

- Opcode `0xc`: **`adc rA, rC`**

  – Mnemonic for when flags not affected: `adc`

  – Mnemonic for when flags affected: `adc.f`

  – Note: Add with Carry

  – Affectable flags: `Z`, `C`, `V`, `N`

- Opcode `0xd`: **`sbc rA, rC`**

  – Mnemonic for when flags not affected: `sbc`

  – Mnemonic for when flags affected: `sbc.f`

  – Note: Subtract with Borrow

  – Affectable flags: `Z`, `C`, `V`, `N`

3.4 Instruction Group 3: Relative Branches

The following encoding is used, with each character representing one bit:

`011i iiii iiii oooo`, where

- `i` is a 9-bit sign-extended immediate, and is denoted `simm`

- `o` is the opcode

Here is a list of instructions from this encoding group.

- Opcode `0x0`: **`bl simm`**

- Name: Branch and Link
- Description: Relative call
- In-Depth Effect: `lr <= pc + 2; pc <= pc + simm + 2;`

- Opcode `0x1`: **`bra`** `simm`

  - Name: BRanch Always
  - Description: Relative branch always
  - Effect: `pc <= pc + simm + 2;`

- Opcode `0x2`: **`beq`** `simm`

  - Name: Branch if EQual
  - Effect: `if (flags.Z) pc <= pc + simm + 2;`

- Opcode `0x3`: **`bne`** `simm`

  - Name: Branch if Not Equal
  - Effect: `if (!flags.Z) pc <= pc + simm + 2;`

- Opcode `0x4`: **`bmi`** `simm`

  - Name: Branch if MInus
  - Effect: `if (flags.N) pc <= pc + simm + 2;`

- Opcode `0x5`: **`bpl`** `simm`

  - Name: Branch if PLus
  - Effect: `if (!flags.N) pc <= pc + simm + 2;`

- Opcode `0x6`: **`bvs`** `simm`

  - Name: Branch if oVerflow Set
  - Effect: `if (flags.V) pc <= pc + simm + 2;`

- Opcode `0x7`: **`bvc`** `simm`

  - Name: Branch if oVerflow Clear
  - Effect: `if (!flags.V) pc <= pc + simm + 2;`

- Opcode `0x8`: **`bgeu`** `simm`

  - Name: Branch if Greater than or Equal Unsigned
  - Effect: `if (flags.C) pc <= pc + simm + 2;`

- Opcode `0x9`: **`bltu`** `simm`

  - Name: Branch if Less Than Unsigned

- Effect: `if (!flags.C) pc <= pc + simm + 2;`

- Opcode `0xa`: **`bgtu simm`**

  - Name: Branch if Greater Than Unsigned
  - Effect: `if (flags.C AND !flags.Z) pc <= pc + simm + 2;`

- Opcode `0xb`: **`bleu simm`**

  - Name: Branch if Less than or Equal Unsigned
  - Effect: `if (!flags.C OR flags.Z) pc <= pc + simm + 2;`

- Opcode `0xc`: **`bges simm`**

  - Name: Branch if Greater than or Equal Signed
  - Effect: `if (!(flags.N XOR flags.V)) pc <= pc + simm + 2;`

- Opcode `0xd`: **`blts simm`**

  - Name: Branch if Less Than Signed
  - Effect: `if (flags.N XOR flags.V) pc <= pc + simm + 2;`

- Opcode `0xe`: **`bgts simm`**

  - Name: Branch if Greater Than Signed
  - Effect: `if (!(flags.N XOR flags.V) AND !flags.Z) pc <= pc + simm + 2;`

- Opcode `0xf`: **`bles simm`**

  - Name: Branch if Less than or Equal Signed
  - Effect: `if ((flags.N XOR flags.V) OR flags.Z) pc <= pc + simm + 2;`

3.5 Instruction Group 4

The following encoding is used, with each character representing one bit:

`100o oooo cccc aaaa`, where

- `o` is the opcode

- `c` encodes register `rC` or register `sC`, where `sC` is one of `hi`, `lo`, `flags`, `ira`, `ids`, or `ie`

- `a` encodes register `rA` or register `sA`, where `sA` is one of `hi`, `lo`, `flags`, `ira`, `ids`, or `ie`

Here is a list of instructions from this encoding group.

- Opcode 0x0: **jl rA**

  – Effect: `lr <= pc + 2; pc <= rA;`

- Opcode 0x1: **jmp rA**

  – Effect: `pc <= rA;`

- Opcode 0x2: **jmp rA, rC**

  – Effect: `pc <= (rA + rC);`

- Opcode 0x3: **jmp ira**

  – Effect: `pc <= ira;`

- Opcode 0x4: **reti**

  – Effect: enables interrupts (by copying 1 into `ie`) and performs `pc <= ira;`

- Opcode 0x5: **cpy rA, sC**

  – Effect: `rA <= sC;`

- Opcode 0x6: **cpy sA, rC**

  – Effect: `sA <= rC;`

- Opcode 0x7: **ei**

  – Effect: copy 1 into `ie`.

- Opcode 0x8: **di**

  – Effect: copy 0 into `ie`.

- Opcode 0x9: **push rA, rC**

  – Effect: pushes `rA` onto the stack, using `rC` as the stack pointer, post-decrementing `rC`.

  – This instruction does nothing when `rA` is the same register as `rC`.

- Opcode 0xa: **pop rA, rC**

  – Effect: pops `rA` off the stack, using `rC` as the stack pointer, pre-incrementing `rC`.

  – This instruction does nothing when `rA` is the same register as `rC`.

- Opcode 0xb: **push sA, rC**

- Effect: pushes `sA` onto the stack, using `rC` as the stack pointer, post-decrementing `rC`.

- Note that `sA` is considered to be 32-bit for the purpose of the store to memory and decrementing `rC`, even if `sA` is `flags` or `ie`.

- Opcode `0xc`: `pop sA, rC`

  - Effect: pops `sA` off the stack, using `rC` as the stack pointer, pre-incrementing `rC`.

  - Note that `sA` is considered to be 32-bit for the purpose of the load from memory and incrementing `rC`, even if `sA` is `flags` or `ie`.

- Opcode `0xd`: `index rA`

  - Effect: Performs `<index_reg> <= rA;` and stores that `index` is in effect.

  - Note: If `index` is in effect and the current instruction is `index`, the current instruction will be treated as a NOP, and `index` will stop being in effect.

  - Note: `pre` and `index` can be combined with one another (though this is only useful for `ldr` and `str`).

  - Note: A non-`pre` instruction following `index` will store that that `index` is not in effect any more. (It will also store that `pre` is not in effect any more).

  - Note: If `index` is in effect, the current instruction cannot be interrupted.

  - Note: Any time `index` stops being in effect, `pre` will stop being in effect as well.

- Opcode `0xe`: `mul rA, rC`

  - Effect: `rA <= rA * rC;`

- Opcode `0xf`: `udiv rA, rC`

  - Effect: `rA <= u32(rA) / u32(rC);`

- Opcode `0x10`: `sdiv rA, rC`

  - Effect: `rA <= s32(rA) / s32(rC);`

- Opcode `0x11`: `umod rA, rC`

  - Effect: `rA <= u32(rA) % u32(rC);`

- Opcode `0x12`: `smod rA, rC`

  - Effect: `rA <= s32(rA) % s32(rC);`

- Opcode 0x13: **lumul rA, rC**

  - Effect: This instruction multiplies **rA** by **rC**, performing an unsigned 32-bit by 32-bit -> 64-bit multiply, storing result in {**hi, lo**}.

- Opcode 0x14: **lsmul rA, rC**

  - Effect: This instruction multiplies **rA** by **rC**, performing a signed 32-bit by 32-bit -> 64-bit multiply, storing result in {**hi, lo**}.

- Opcode 0x15: **ludiv rA, rC**

  - Effect: performs a 64-bit by 64-bit unsigned division of {**hi, lo**} by {**rA, rC**}, storing 64-bit result in {**hi, lo**}.
  - Note: This instruction executes more quickly if **rA**'s value is 0x00000000, i.e. if the operation is actually a 64-bit by 32-bit -> 64-bit unsigned divide.

- Opcode 0x16: **lsdiv rA, rC**

  - Effect: performs a 64-bit by 64-bit signed division of {**hi, lo**} by {**rA, rC**}, storing 64-bit result in {**hi, lo**}.
  - Note: This instruction executes more quickly if **rA** is equal to bits [63:32] of **sign_extend_to_64(rC)**.

- Opcode 0x17: **lumod rA, rC**

  - Effect: performs a 64-bit by 64-bit unsigned modulo of {**hi, lo**} by {**rA, rC**}, storing 64-bit result in {**hi, lo**}.
  - Note: This instruction executes more quickly if **rA**'s value is 0x00000000, i.e. if the operation is actually a 64-bit by 32-bit -> 64-bit unsigned modulo.

- Opcode 0x18: **lsmod rA, rC**

  - Effect: performs a 64-bit by 64-bit signed modulo of {**hi, lo**} by {**rA, rC**}, storing 64-bit result in {**hi, lo**}.
  - Note: This instruction executes more quickly if **rA** is equal to bits [63:32] of **sign_extend_to_64(rC)**.

- Opcode 0x19: **ldub rA, [rC]**

  - Effect: Load an 8-bit value from memory at address computed as **rC** + **<index_reg>**, zero-extend the 8-bit value to 32 bits, then put the zero-extended 32-bit value into **rA**.
  - The **<index_reg>** value is guaranteed to be zero unless an **index** is in effect.
  - Shorthand for having the assembler insert an **index rB** instruction before this one: **ldub rA, [rC, rB]**

- Opcode `0x1a`: **`ldsb rA, [rC]`**

  - Effect: Load an 8-bit value from memory at address computed as `rC` `+ <index_reg>`, sign-extend the 8-bit value to 32 bits, then put the sign-extended 32-bit value into `rA`.
  - The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.
  - Shorthand for having the assembler insert an `index rB` instruction before this one: `ldsb rA, [rC, rB]`

- Opcode `0x1b`: **`lduh rA, [rC]`**

  - Effect: Load a 16-bit value from memory at address computed as `rC` `+ <index_reg>`, zero-extend the 16-bit value to 32 bits, then put the zero-extended 32-bit value into `rA`.
  - The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.
  - Shorthand for having the assembler insert an `index rB` instruction before this one: `lduh rA, [rC, rB]`

- Opcode `0x1c`: **`ldsh rA, [rC]`**

  - Effect: Load a 16-bit value from memory at address computed as `rC` `+ <index_reg>`, sign-extend the 16-bit value to 32 bits, then put the zero-extended 32-bit value into `rA`.
  - The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.
  - Shorthand for having the assembler insert an `index rB` instruction before this one: `ldsh rA, [rC, rB]`

- Opcode `0x1d`: **`stb rA, [rC]`**

  - Effect: Store `rA[7:0]` to memory at the address computed as `rC +` `<index_reg>`.
  - The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.
  - Shorthand for having the assembler insert an `index rB` instruction before this one: `stb rA, [rC, rB]`

- Opcode `0x1e`: **`sth rA, [rC]`**

  - Effect: Store `rA[15:0]` to memory at the address computed as `rC +` `<index_reg>`.
  - The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.
  - Shorthand for having the assembler insert an `index rB` instruction before this one: `sth rA, [rC, rB]`

3.6   Instruction Group 5: Immediate Indexed Load

The following encoding is used, with each character representing one bit:

`101i iiii cccc aaaa`, where

- `i` is a 5-bit sign-extended immediate, which can be expanded by `pre`, and is denoted `simm`
- `c` encodes register `rC`
- `a` encodes register `rA`

The one instruction from this encoding group is `ldr rA, [rC, #simm]`. This is a 32-bit load into `rA`, where the effective address to load from is computed as `rC + <index_reg> + simm`, using the sign-extended form of `simm`.

The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.

Shorthand for having the assembler insert an `index rB` instruction before this one: `ldr rA, [rC, rB, #simm]`

3.7   Instruction Group 6: Immediate Indexed Store

The following encoding is used, with each character representing one bit:

`110i iiii cccc aaaa`, where

- `i` is a 5-bit sign-extended immediate, which can be expanded by `pre`
- `c` encodes register `rC`
- `a` encodes register `rA`

The one instruction from this encoding group is `str rA, [rC, #simm]`. This is a 32-bit store of `rA`, where the effective address to store to is computed as `<index_reg> + rC + simm`, using the sign-extended form of `simm`.

The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.

Shorthand for having the assembler insert an `index rB` instruction before this one: `str rA, [rC, rB, #simm]`