

Flare32 CPU

FL4SHK

May 15, 2023

Table of Contents

Table of Contents	1
Introduction	2
2.1 General Information	2
2.2 Registers	2
Instruction Set	4
3.1 Instruction Group 0: pre and lpre	4
3.1.1 Handling of pre , lpre , and index	5
3.2 Instruction Group 1	6
3.3 Instruction Group 2	8
3.4 Instruction Group 3: Relative Branches	11
3.5 Instruction Group 4	13
3.6 Instruction Group 5: Immediate Indexed Load	18
3.7 Instruction Group 6: Immediate Indexed Store	18
3.8 Instruction Group 7, Subgroup 0b00: Extra 8-bit and 16-bit Ops	19
3.9 Instruction Group 7, Subgroup 0b010: Extra load/store instructions	20

Introduction

2.1 General Information

- Addresses are 32-bit.
- Big-endian byte ordering is used.
- Bytes are octets (8 bits).
- Instructions must be aligned to 16 bits, so jump and branch targets must also be aligned to 16 bits.
 - Branch offsets encoded into instructions must be 16-bit aligned, or in other words bit 0 of the branch offset must be 0b0

2.2 Registers

There are sixteen general-purpose registers (all of which are 32-bit): `r0`, `r1`, `r2`, ..., `r11`, `r12`, `lr` (link register, the return address of `bl` and `jl`), `fp` (frame pointer), `sp` (stack pointer).

The program counter, `pc`, is 32 bits long, as addresses are 32-bit.
Here are the special-purpose registers:

- `flags`: arithmetic/logic FLAGS; (reg encoding: 0x0)
- `ids`: Interrupt DeStination: the address to jump to upon an interrupt being serviced (also known as the interrupt vector); (reg encoding: 0x1)
- `ira`: Interrupt Return Address: the address that was jumped from upon an interrupt being serviced; (reg encoding: 0x2)
- `ie`: IRQ Enable flag: flag indicating whether IRQs are disabled (0x0) or enabled (0x1); note that this flag starts with a value of 0x0; (reg encoding: 0x3)
- `ity`: Interrupt TYpe: flag indicating whether the most recently taken interrupt is an IRQ (0x0) or a `swi` (0x1) (reg encoding: 0x4)
- `sty`: Software interrupt TYpe: `swi`'s argument. For `swi rA, #simm`, this is the value `rA + simm` (reg encoding: 0x5)
- Note: All other encodings for special-purpose registers are reserved.

Here are the bits of `flags`:

- Zero (Z): (`flags` bit 0)

- Carry (C): (**flags** bit 1)
- oVerflow (V): (**flags** bit 2)
- Negative (N): (**flags** bit 3)
- Note: All other bit positions of **flags** are reserved.

Instruction Set

3.1 Instruction Group 0: **pre** and **lpre**

For **pre**, the following encoding is used, with each character representing one bit:

0000 *iiii* *iiii* *iiii*, where

- *i* is a 12-bit constant.

For **lpre**, the following encoding is used, with each character representing one bit:

0001 0*iiii* *iiii* *iiii* *iiii* *iiii* *iiii* *iiii*, where

- *i* is a 27-bit constant.

pre and **lpre** are mechanisms by which immediates larger than normal can be used, essentially acting like variable width instructions.

There is no mechanism in the assembly language itself to use **pre** or **lpre** as instructions. Instead, it is expected that the assembler or linker will be the one to insert **pre** or **lpre** as needed if an immediate is too large for a particular instruction.

For non-branch instructions:

- If **pre** is used, the immediate field of the **pre** instruction will form bits [16:5] of the immediate of the next non-**index** instruction. The 17-bit immediate will then be sign-extended to 32 bits.
- On the other hand, if **lpre** is used, the immediate field of the **lpre** instruction will form bits [31:5] of the immediate of the next non-**index** instruction.

For branch instructions (group 3):

- if **pre** is used, the immediate field of the **pre** instruction will form bits [20:9] of the immediate of the next non-**index** instruction. The 21-bit immediate will then be sign-extended to 32 bits.
- On the other hand, if **lpre** is used, the immediate field of the **lpre** instruction will form bits [31:9] of the immediate of the next non-**index** instruction.

Instructions that use an immediate value, other than **pre**, **lpre**, and relative branches normally have 5-bit immediates. Relative branches normally have 9-bit branch offsets. if **pre** or **lpre** is used with a relative branch, bits [9:5] of the branch offset that is encoded directly into a relative branch instruction will be ignored, and the immediate field of the **pre** or **lpre** instruction will be used to determine bits [9:5] of the branch offset..

3.1.1 Handling of pre, lpre, and index

When a **pre** or **lpre** instruction is found, **pre** or **lpre** will be considered to be "in effect". This condition lasts for one or two instructions after the **pre** or **lpre** instruction, depending on whether or not **index** was in effect.

index is an instruction (defined later) that allows a subsequent load or store instruction to perform **base_reg + index_reg** indexing. When an **index** instruction is found, it is considered to be in effect. Like **pre** and **lpre**, **index** is inserted automatically by the assembler.

index can be combined with **pre** or **lpre**, and it does not matter whether **index** or the **pre/lpre** instruction came first.

When **pre**, **lpre**, or **index** is in effect, IRQs will not be serviced.

Pseudo code for handling the how **pre**, **lpre**, and **index** are to be handled regarding whether or not they are "in effect" is as follows:

```
function handle_pre_lpre_index(input in, output out) {
    if (in.instruction.is_pre()) {
        if (in.state.pre.have || in.state.lpre.have) {
            out.state.can_service_interrupts = false;

            // Instruction was a NOP
            out.state.pre.have = false;
            out.state.lpre.have = false;
            out.state.index.have = false;
            out.instruction = nop;
        } else {
            if (in.state.index.have) {
                out.state.can_service_interrupts = false;
            } else {
                out.state.can_service_interrupts = true;
            }
            out.state.pre.set_have(true);
            out.instruction = in.instruction;
        }
    } else if (in.instruction.is_lpre()) {
        if (in.state.pre.have || in.state.lpre.have) {
            out.state.can_service_interrupts = false;

            // Instruction was a NOP
            out.state.pre.have = false;
            out.state.lpre.have = false;
            out.state.index.have = false;
            out.instruction = nop;
        } else {
            if (in.state.index.have) {
                out.state.can_service_interrupts = false;
            }
        }
    }
}
```

```

        } else {
            out.state.can_service_interrupts = true;
        }
        out.state.lpre.set_have(true);
        out.instruction = in.instruction;
    }
} else if (in.instruction.is_index()) {
    if (in.state.index.have) {
        out.state.can_service_interrupts = false;

        // Instruction was a NOP
        out.state.pre.have = false;
        out.state.lpre.have = false;
        out.state.index.have = false;
        out.instruction = nop;
    } else {
        if (in.state.pre.have || in.state.lpre.have) {
            out.state.can_service_interrupts = false;
        } else {
            out.state.can_service_interrupts = true;
        }
        out.state.index.have = true;
        out.instruction = in.instruction;
    }
} else {
    if (
        in.state.pre.have || in.state.lpre.have || in.state.index.have
    ) {
        out.state.can_service_interrupts = false;
    } else {
        out.state.can_service_interrupts = true;
    }
    out.instruction = instruction;

    // Whenever we see an instruction other than pre, lpre, or index,
    // that means that those instructions stop being "in effect".
    out.state.pre.have = false;
    out.state.lpre.have = false;
    out.state.index.have = false;
}
}

```

3.2 Instruction Group 1

The following encoding is used, with each character representing one bit:
001i iiii oooo aaaa, where

- `i` is a 5-bit sign-extended immediate, and is denoted `simm`
- `a` encodes register `rA`
- `o` is the opcode

Here is a list of instructions from this encoding group.

- Opcode 0x0: `add rA, #simm`
- Opcode 0x1: `add rA, pc, #simm`
 - Effect: `rA <= pc + simm + 2;`
- Opcode 0x2: `add rA, sp, #simm`
- Opcode 0x3: `add rA, fp, #simm`
- Opcode 0x4: `cmp rA, #simm`
 - Effect: Compare `rA` to `simm`.
 - Affectable flags: Z, C, V, N
- Opcode 0x5: `cpy rA, #simm`
 - Effect: Copy an immediate value into `rA`
- Opcode 0x6: `lsl rA, #simm`
 - Effect: Logical shift left
- Opcode 0x7: `lsr rA, #simm`
 - Effect: Logical shift right
- Opcode 0x8: `asr rA, #simm`
 - Effect: Arithmetic shift right
- Opcode 0x9: `and rA, #simm`
 - Effect: Bitwise AND
- Opcode 0xa: `orr rA, #simm`
 - Effect: Bitwise OR
- Opcode 0xb: `xor rA, #simm`
 - Effect: Bitwise XOR
- Opcode 0xc: `ze rA, #simm`

- Effect: Set `rA[31:simm]` to zero.
- Opcode 0xd: `se rA, #simm`
 - Effect: Sign-extend `rA[simm:0]` to 32 bits, then copy that value to `rA`.
- Opcode 0xe: `swi rA, #simm`
 - Effect: Call software interrupt number `rA + simm`.
- Opcode 0xf: `swi #simm`
 - Effect: Call software interrupt number `simm`.

3.3 Instruction Group 2

The following encoding is used, with each character representing one bit:

010f oooo bbbb aaaa, where

- o is the opcode
- b encodes register `rB`
- a encodes register `rA`
- f is encoded as 0 if this instruction cannot affect flags and encoded 1 if this instruction is permitted to affect flags. Note that `cmp` is permitted to affect flags regardless of this bit.

Here is a list of instructions from this encoding group.

- Opcode 0x0: `add rA, rB`
 - Mnemonic for when flags not affected: `add`
 - Mnemonic for when flags affected: `add.f`
 - Affectable flags: Z, C, V, N
- Opcode 0x1: `sub rA, rB`
 - Mnemonic for when flags not affected: `sub`
 - Mnemonic for when flags affected: `sub.f`
 - Affectable flags: Z, C, V, N
- Opcode 0x2: `add rA, sp, rB`
 - Mnemonic for when flags not affected: `add`

- Mnemonic for when flags affected: **add.f**
 - Affectable flags: Z, C, V, N
- Opcode 0x3: **add rA, fp, rB**
 - Mnemonic for when flags not affected: **add**
 - Mnemonic for when flags affected: **add.f**
 - Affectable flags: Z, C, V, N
- Opcode 0x4: **cmp rA, rB**
 - Effect: Compare **rA** to **rB**. **cmp** is *always* able to affect flags, independent of the encoded **f** bit of the instruction.
 - Affectable flags: Z, C, V, N
- Opcode 0x5: **cpy rA, rB**
 - Mnemonic for when flags not affected: **cpy**
 - Mnemonic for when flags affected: **cpy.f**
 - Effect: Copy **rB** into **rA**
 - Affectable flags: Z, N
- Opcode 0x6: **lsl rA, rB**
 - Mnemonic for when flags not affected: **lsl**
 - Mnemonic for when flags affected: **lsl.f**
 - Effect: Logical shift left
 - Affectable flags: Z, N
- Opcode 0x7: **lsr rA, rB**
 - Mnemonic for when flags not affected: **lsr**
 - Mnemonic for when flags affected: **lsr.f**
 - Effect: Logical shift right
 - Affectable flags: Z, N
- Opcode 0x8: **asr rA, rB**
 - Mnemonic for when flags not affected: **asr**
 - Mnemonic for when flags affected: **asr.f**
 - Effect: Arithmetic shift right
 - Affectable flags: Z, N
- Opcode 0x9: **and rA, rB**
 - Mnemonic for when flags not affected: **and**

- Mnemonic for when flags affected: **and.f**
- Effect: Bitwise AND
- Affectable flags: Z, N
- Opcode 0xa: **orr rA, rB**
 - Mnemonic for when flags not affected: **orr**
 - Mnemonic for when flags affected: **orr.f**
 - Effect: Bitwise OR
 - Affectable flags: Z, N
- Opcode 0xb: **xor rA, rB**
 - Mnemonic for when flags not affected: **xor**
 - Mnemonic for when flags affected: **xor.f**
 - Effect: Bitwise XOR
 - Affectable flags: Z, N
- Opcode 0xc: **adc rA, rB**
 - Mnemonic for when flags not affected: **adc**
 - Mnemonic for when flags affected: **adc.f**
 - Effect: Add with Carry, using the formula $rA + rB + \text{flags.C}$ to compute the value that will be written into rA.
 - Affectable flags: Z, C, V, N
- Opcode 0xd: **sbc rA, rB**
 - Mnemonic for when flags not affected: **sbc**
 - Mnemonic for when flags affected: **sbc.f**
 - Effect: Subtract with Borrow, using the formula $rA + (rB) + \text{flags.C}$ to compute the value that will be written into rA.
 - Affectable flags: Z, C, V, N
- Opcode 0xe: **cmpbc rA, rB**
 - Effect: Compare rA to rB, but with carry-in and a different effect for setting the flags.Z. **cmpbc** is *always* able to affect flags, independent of the encoded f bit of the instruction.
 - Note: this instruction acts much like **sbc rA, rB**, but without storing the subtraction's result into rA. However, this instruction sets the Z flag to $\text{prev}(\text{flags.Z}) \text{ AND } ((rA + (rB) + \text{flags.C}) == 0)$
 - Affectable flags: Z, C, V, N

3.4 Instruction Group 3: Relative Branches

The following encoding is used, with each character representing one bit:

011i iiii iiii oooo, where

- i is a 9-bit sign-extended immediate, and is denoted **simm**
- o is the opcode

Here is a list of instructions from this encoding group.

- Opcode 0x0: **bl** **simm**
 - Name: Branch and Link
 - Description: Relative call
 - Effect: $lr \leq pc + 2$; $pc \leq pc + simm + 2$;
- Opcode 0x1: **bra** **simm**
 - Name: BRanch Always
 - Description: Unconditional relative branch
 - Effect: $pc \leq pc + simm + 2$;
- Opcode 0x2: **beq** **simm**
 - Name: Branch if EQual
 - Effect: if (flags.Z) $pc \leq pc + simm + 2$;
- Opcode 0x3: **bne** **simm**
 - Name: Branch if Not Equal
 - Effect: if (!flags.Z) $pc \leq pc + simm + 2$;
- Opcode 0x4: **bmi** **simm**
 - Name: Branch if MInus
 - Effect: if (flags.N) $pc \leq pc + simm + 2$;
- Opcode 0x5: **bpl** **simm**
 - Name: Branch if PLus
 - Effect: if (!flags.N) $pc \leq pc + simm + 2$;
- Opcode 0x6: **bvs** **simm**
 - Name: Branch if oVerflow Set
 - Effect: if (flags.V) $pc \leq pc + simm + 2$;

- Opcode 0x7: **bvc** *simm*
 - Name: Branch if oVerflow Clear
 - Effect: if (**!flags.V**) $pc \leq pc + \text{simm} + 2$;
- Opcode 0x8: **bgeu** *simm*
 - Name: Branch if Greater than or Equal Unsigned
 - Effect: if (**flags.C**) $pc \leq pc + \text{simm} + 2$;
- Opcode 0x9: **bltu** *simm*
 - Name: Branch if Less Than Unsigned
 - Effect: if (**!flags.C**) $pc \leq pc + \text{simm} + 2$;
- Opcode 0xa: **bgtu** *simm*
 - Name: Branch if Greater Than Unsigned
 - Effect: if (**flags.C AND !flags.Z**) $pc \leq pc + \text{simm} + 2$;
- Opcode 0xb: **bleu** *simm*
 - Name: Branch if Less than or Equal Unsigned
 - Effect: if (**!flags.C OR flags.Z**) $pc \leq pc + \text{simm} + 2$;
- Opcode 0xc: **bges** *simm*
 - Name: Branch if Greater than or Equal Signed
 - Effect: if (**!(flags.N XOR flags.V)**) $pc \leq pc + \text{simm} + 2$;
- Opcode 0xd: **blts** *simm*
 - Name: Branch if Less Than Signed
 - Effect: if (**flags.N XOR flags.V**) $pc \leq pc + \text{simm} + 2$;
- Opcode 0xe: **bgts** *simm*
 - Name: Branch if Greater Than Signed
 - Effect: if (**!(flags.N XOR flags.V) AND !flags.Z**) $pc \leq pc + \text{simm} + 2$;
- Opcode 0xf: **bles** *simm*
 - Name: Branch if Less than or Equal Signed
 - Effect: if (**(flags.N XOR flags.V) OR flags.Z**) $pc \leq pc + \text{simm} + 2$;

3.5 Instruction Group 4

The following encoding is used, with each character representing one bit:

100o oooo bbbb aaaa, where

- o is the opcode
- b encodes register **rB** or register **sC**, where **sC** is one of the special-purpose registers
- a encodes register **rA** or register **sA**, where **sA** is one of the special-purpose registers

Here is a list of instructions from this encoding group.

- Opcode 0x0: **j1 rA**
 - Effect: $lr \leq pc + 2$; $pc \leq rA$;
- Opcode 0x1: **jmp rA**
 - Effect: $pc \leq rA$;
- Opcode 0x2: **jmp ira**
 - Effect: $pc \leq ira$;
- Opcode 0x3: **reti**
 - Effect: enables IRQs (by copying 0x1 into **ie**) and performs $pc \leq ira$;
- Opcode 0x4: **ei**
 - Effect: copy 1 into **ie**.
- Opcode 0x5: **di**
 - Effect: copy 0 into **ie**.
- Opcode 0x6: **push rA, rB**
 - Effect: pushes **rA** onto the stack, using **rB** as the stack pointer, post-decrementing **rB**.
 - This instruction does nothing when **rA** is the same register as **rB**.
 - Note: As a pseudo instruction, omitting ", **rB**" will automatically select **sp** as the particular stack pointer. **sp**
- Opcode 0x7: **push sA, rB**

- Effect: pushes **sA** onto the stack, using **rB** as the stack pointer, post-decrementing **rB**.
- Note that **sA** is considered to be 32-bit for the purpose of the store to memory and decrementing **rB**, even if **sA** is **flags** or **ie**.
- Note: As a pseudo instruction, omitting ", **rB**" will automatically select **sp** as the particular stack pointer.
- Opcode 0x8: **pop rA, rB**
 - Effect: pops **rA** off the stack, using **rB** as the stack pointer, pre-incrementing **rB**.
 - This instruction does nothing when **rA** is the same register as **rB**.
 - Note: As a pseudo instruction, omitting ", **rB**" will automatically select **sp** as the particular stack pointer.
- Opcode 0x9: **pop sA, rB**
 - Pseudo instruction with same effect:
 - * **ldrib sA, rB**
 - * Name explanation: load register, increment base before
 - Effect: pops **sA** off the stack, using **rB** as the stack pointer, pre-incrementing **rB**.
 - Note that **sA** is considered to be 32-bit for the purpose of the load from memory and incrementing **rB**, even if **sA** is **flags** or **ie**.
 - Note: As a pseudo instruction, omitting ", **rB**" will automatically select **sp** as the particular stack pointer.
- Opcode 0xa: **index rA**
 - Effect: Performs **<index_reg> <= rA;** and stores that **index** is in effect.
 - Note: If **index** is in effect and the current instruction is **index**, the current instruction will be treated as a NOP, and **index** will stop being in effect.
 - Note: **pre** and **index** can be combined with one another (though this is only useful for **ldr** and **str**).
 - Note: A non-**pre** instruction following **index** will store that that **index** is not in effect any more. (It will also store that **pre** is not in effect any more).
 - Note: If **index** is in effect, the current instruction cannot be interrupted by an IRQ.
 - Note: Any time **index** stops being in effect, **pre** will stop being in effect as well.

- Opcode 0xb: **mul** *rA*, *rB*
 - Effect: $rA \leftarrow rA * rB$;
- Opcode 0xc: **udiv** *rA*, *rB*
 - Effect: $rA \leftarrow u32(rA) / u32(rB)$;
- Opcode 0xd: **sdiv** *rA*, *rB*
 - Effect: $rA \leftarrow s32(rA) / s32(rB)$;
- Opcode 0xe: **umod** *rA*, *rB*
 - Effect: $rA \leftarrow u32(rA) \% u32(rB)$;
- Opcode 0xf: **smod** *rA*, *rB*
 - Effect: $rA \leftarrow s32(rA) \% s32(rB)$;
- Opcode 0x10: **lumul** *rA*, *rB*
 - Effect: This instruction multiplies *rA* by *rB*, performing an unsigned 32-bit by 32-bit \rightarrow 64-bit multiply, storing result in `concat{r0, r1}`.
- Opcode 0x11: **lsmul** *rA*, *rB*
 - Effect: This instruction multiplies *rA* by *rB*, performing a signed 32-bit by 32-bit \rightarrow 64-bit multiply, storing result in `concat{r0, r1}`.
- Opcode 0x12: **ludiv** *rA*, *rB*
 - Effect: performs a 64-bit by 64-bit unsigned division of `concat{rA, r{A + 1}}` by `concat{rB, r{B + 1}}`, storing 64-bit result in `concat{rA, r{A + 1}}`.
 - Note: This instruction operates as if *rA* were encoded with ($A[0x0] == 0b0$) and *rB* were encoded with ($B[0x0] == 0b0$),
 - Note: This instruction executes more quickly if *rB*'s value is 0x00000000, i.e. if the operation is actually a 64-bit by 32-bit \rightarrow 64-bit unsigned divide.
- Opcode 0x13: **lsdiv** *rA*, *rB*
 - Effect: performs a 64-bit by 64-bit signed division of `concat{rA, r{A + 1}}` by `concat{rB, r{B + 1}}`, storing 64-bit result in `concat{rA, r{A + 1}}`.
 - Note: This instruction operates as if *rA* were encoded with ($A[0x0] == 0b0$) and *rB* were encoded with ($B[0x0] == 0b0$),
 - Note: This instruction executes more quickly if *rB* is equal to bits [63:32] of $s64(r\{B+1\})$.

- Opcode 0x14: **lumod** rA, rB
 - Effect: performs a 64-bit by 64-bit unsigned modulo of `concat{rA, r{A + 1}}` by `concat{rB, r{B + 1}}`, storing 64-bit result in `concat{rA, r{A + 1}}`.
 - Note: This instruction operates as if rA were encoded with (A[0x0] == 0b0) and rB were encoded with (B[0x0] == 0b0),
 - Note: This instruction executes more quickly if rB's value is 0x00000000, i.e. if the operation is actually a 64-bit by 32-bit -> 64-bit unsigned modulo.
- Opcode 0x15: **lsmod** rA, rB
 - Effect: performs a 64-bit by 64-bit signed modulo of `concat{rA, r{A + 1}}` by `concat{rB, r{B + 1}}`, storing 64-bit result in `concat{rA, r{A + 1}}`.
 - Note: This instruction operates as if rA were encoded with (A[0x0] == 0b0) and rB were encoded with (B[0x0] == 0b0),
 - Note: This instruction executes more quickly if rB is equal to bits [63:32] of `s64(r{B+1})`.
- Opcode 0x16: **ldub** rA, [rB]
 - Effect: Load an 8-bit value from memory at address computed as `rB + <index_reg>`, zero-extend the 8-bit value to 32 bits, then put the zero-extended 32-bit value into rA.
 - The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.
 - Shorthand for having the assembler insert an `index rC` instruction before this one: `ldub rA, [rB, rC]`
- Opcode 0x17: **ldsb** rA, [rB]
 - Effect: Load an 8-bit value from memory at address computed as `rB + <index_reg>`, sign-extend the 8-bit value to 32 bits, then put the sign-extended 32-bit value into rA.
 - The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.
 - Shorthand for having the assembler insert an `index rC` instruction before this one: `ldsb rA, [rB, rC]`
- Opcode 0x18: **lduh** rA, [rB]
 - Effect: Load a 16-bit value from memory at address computed as `rB + <index_reg>`, zero-extend the 16-bit value to 32 bits, then put the zero-extended 32-bit value into rA.

- The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.
 - Shorthand for having the assembler insert an `index rC` instruction before this one: `lduh rA, [rB, rC]`
- Opcode 0x19: `ldsh rA, [rB]`
 - Effect: Load a 16-bit value from memory at address computed as `rB + <index_reg>`, sign-extend the 16-bit value to 32 bits, then put the zero-extended 32-bit value into `rA`.
 - The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.
 - Shorthand for having the assembler insert an `index rC` instruction before this one: `ldsh rA, [rB, rC]`
- Opcode 0x1a: `stb rA, [rB]`
 - Effect: Store `rA[7:0]` to memory at the address computed as `rB + <index_reg>`.
 - The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.
 - Shorthand for having the assembler insert an `index rC` instruction before this one: `stb rA, [rB, rC]`
- Opcode 0x1b: `sth rA, [rB]`
 - Effect: Store `rA[15:0]` to memory at the address computed as `rB + <index_reg>`.
 - The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.
 - Shorthand for having the assembler insert an `index rC` instruction before this one: `sth rA, [rB, rC]`
- Opcode 0x1c: `cpy rA, sC`
 - Effect: `rA <= sC;`
- Opcode 0x1d: `cpy sA, rB`
 - Effect: `sA <= rB;`
- Opcode 0x1e: `cpy sA, sB`
 - Effect: `sA <= sB;`

3.6 Instruction Group 5: Immediate Indexed Load

The following encoding is used, with each character representing one bit:

101i iiii bbbb aaaa, where

- i is a 5-bit sign-extended immediate, which can be expanded by `pre`, and is denoted `simm`
- b encodes register `rB`
- a encodes register `rA`

The one instruction from this encoding group is `ldr rA, [rB, #simm]`. This is a 32-bit load into `rA`, where the effective address to load from is computed as `rB + <index_reg> + simm`, using the sign-extended form of `simm`.

The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.

Shorthand for having the assembler insert an `index rC` instruction before this one: `ldr rA, [rB, rC, #simm]`

3.7 Instruction Group 6: Immediate Indexed Store

The following encoding is used, with each character representing one bit:

110i iiii bbbb aaaa, where

- i is a 5-bit sign-extended immediate, which can be expanded by `pre`
- b encodes register `rB`
- a encodes register `rA`

The one instruction from this encoding group is `str rA, [rB, #simm]`. This is a 32-bit store of `rA`, where the effective address to store to is computed as `<index_reg> + rB + simm`, using the sign-extended form of `simm`.

The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.

Shorthand for having the assembler insert an `index rC` instruction before this one: `str rA, [rB, rC, #simm]`

3.8 Instruction Group 7, Subgroup 0b00: Extra 8-bit and 16-bit Ops

The following encoding is used, with each character representing one bit:

1110 0w0o bbbb aaaa, where

- **w** is the operation width
 - When 0b0: 8-bit operation
 - When 0b1: 16-bit operation
- **o** is the opcode
- **b** encodes register **rB**
- **a** encodes register **rA**

Here is a list of instructions from this encoding group.

- Opcode 0x0:
 - **w** value 0b0: **cmpb rA, rB**
 - * Effect: Compare **rA[7:0]** to **rB[7:0]**
 - * Affectable flags: Z, C, V, N
 - **w** value 0b1: **cmph rA, rB**
 - * Effect: Compare **rA[15:0]** to **rB[15:0]**
 - * Affectable flags: Z, C, V, N
- Opcode 0x1:
 - **w** value 0b0: **lsrb rA, rB**
 - * Effect: Logical shift right **rA[7:0]** by **rB**
 - **w** value 0b1: **lsrh rA, rB**
 - * Effect: Logical shift right **rA[15:0]** by **rB**
- Opcode 0x2:
 - **w** value 0b0: **asrb rA, rB**
 - * Effect: Arithmetic shift right **rA[7:0]** by **rB**
 - **w** value 0b1: **asrh rA, rB**
 - * Effect: Arithmetic shift right **rA[15:0]** by **rB**

3.9 Instruction Group 7, Subgroup 0b010: Extra load/store instructions

The following encoding is used, with each character representing one bit:

1110 10oo bbbb aaaa, where

- o is the opcode
- b encodes register **sA**
- a encodes register **rB** or **sB**

Here is a list of instructions from this encoding group.

- Opcode 0x0: **ldr sA, [rB]**
 - Effect: 32-bit load of **sA** from memory at address held in **rB**.
- Opcode 0x1: **ldr sA, [sB]**
 - Effect: 32-bit load of **sA** from memory at address held in **sB**.
- Opcode 0x2: **str sA, [rB]**
 - Effect: 32-bit store of **sA** to memory at address held in **rB**.
- Opcode 0x3: **str sA, [sB]**
 - Effect: 32-bit store of **sA** to memory at address held in **sB**.