1. Introduction

1.1. General Information
- Addresses are 32-bit.
- Big-endian byte ordering is used.
- Bytes are octets (8 bits).
- Instructions must be aligned to 16 bits, so jump and branch targets must also be aligned to 16 bits.
  - Branch offsets encoded into instructions must be 16-bit aligned, or in other words bit 0 of the branch offset must be `0b0`

1.2. Registers
There are sixteen general-purpose registers (all of which are 32-bit): `r0`, `r1`, `r2`, ..., `r11`, `r12`, `lr` (link register, the return address of `bl` and `jl`), `fp` (frame pointer), and `sp` (stack pointer).

The program counter, `pc`, is 32 bits long, as addresses are 32-bit.

Here are the special-purpose registers:
- `flags`: arithmetic/logic FLAGS; (reg encoding: `0x0`)
- `ids`: Interrupt DeStination: the address to jump to upon an interrupt being serviced (also known as the interrupt vector); (reg encoding: `0x1`)
- `ira`: Interrupt Return Address: the address that was jumped from upon an interrupt being serviced; (reg encoding: `0x2`)
- `ie`: IRQ Enable flag: flag indicating whether IRQs are disabled (`0x0`) or enabled (`0x1`); note that this flag starts with a value of `0x0`; (reg encoding: `0x3`)
- `ity`: Interrupt TYpe: flag indicating whether the most recently taken interrupt is an IRQ (`0x0`) or a `swi` (`0x1`) (reg encoding: `0x4`)
- `sty`: Software interrupt TYpe: `swi`'s argument. For `swi rA, #simm`, this is the value `rA + simm` (reg encoding: `0x5`)
- Note: All other encodings for special-purpose registers are reserved.

Here are the bits of `flags`:
- Zero (Z): (`flags` bit 0)
- Carry (C): (`flags` bit 1)
- oVerflow (V): (`flags` bit 2)
- Negative (N): (`flags` bit 3)
- Note: All other bit positions of `flags` are reserved.

2. Instruction Set

2.1. Instruction Group 0: `pre` and `lpre`
For `pre`, the following encoding is used, with each character representing one bit:
`0000 iiii iiii iiii`, where
- `i` is a 12-bit constant.

For `lpre`, the following encoding is used, with each character representing one bit:
`0001 0iii iiii iiii  iiii iiii iiii iiii`, where
- `i` is a 27-bit constant.

`pre` and `lpre` are mechanisms by which immediates larger than normal can be used, essentially acting like variable width instructions.

There is no mechanism in the assembly language itself to use `pre` or `lpre` as instructions. Instead, it is expected that the assembler or linker will be the one to insert `pre` or `lpre` as needed if an immediate is too large for a particular instruction.

For non-branch instructions:

- If `pre` is used, the immediate field of the `pre` instruction will form bits `[16:5]` of the immediate of the next non-`index` instruction. The 17-bit immediate will then be sign-extended to 32 bits. Sign-extension will be performed on the 17-bit immediate even if, had there been no `pre`, the 5-bit immediate would have been zero-extended.
- On the other hand, if `lpre` is used, the immediate field of the `lpre` instruction will form bits `[31:5]` of the immediate of the next non-`index` instruction.

For branch instructions (group 3):
- if `pre` is used, the immediate field of the `pre` instruction will form bits `[20:9]` of the immediate of the next non-`index` instruction. The 21-bit immediate will then be sign-extended to 32 bits.
- On the other hand, if `lpre` is used, the immediate field of the `lpre` instruction will form bits `[31:9]` of the immediate of the next non-`index` instruction.

2.1.1. Handling of `pre`, `lpre`, and `index`

When a `pre` or `lpre` instruction is found, `pre` or `lpre` will be considered to be "in effect". This condition lasts for one or two instructions after the `pre` or `lpre` instruction, depending on whether or not `index` was in effect.

`index` is an instruction (defined later) that allows a subsequent load or store instruction to perform `base_reg + index_reg` indexing. When an `index` instruction is found, it is considered to be in effect. Like `pre` and `lpre`, `index` is inserted automatically by the assembler.

`index` can be combined with `pre` or `lpre`, and it does not matter whether `index` or the `pre`/`lpre` instruction came first.

When `pre`, `lpre`, or `index` is in effect, IRQs will not be serviced.

Pseudo code for handling the how `pre`, `lpre`, and `index` are to be handled regarding whether or not they are "in effect" is as follows:

```
function handle_pre_lpre_index(input in, output out) {
  if (in.instruction.is_pre()) {
    if (in.state.pre.have || in.state.lpre.have) {
      out.state.can_service_interrupts = false;

      // Instruction was a NOP
      out.state.pre.have = false;
      out.state.lpre.have = false;
      out.state.index.have = false;
      out.instruction = nop;
    } else {
      if (in.state.index.have) {
        out.state.can_service_interrupts = false;
      } else {
        out.state.can_service_interrupts = true;
      }
      out.state.pre.set_have(true);
      out.instruction = in.instruction;
    }
  } else if (in.instruction.is_lpre()) {
    if (in.state.pre.have || in.state.lpre.have) {
      out.state.can_service_interrupts = false;

      // Instruction was a NOP
      out.state.pre.have = false;
      out.state.lpre.have = false;
      out.state.index.have = false;
      out.instruction = nop;
    } else {
      if (in.state.index.have) {
        out.state.can_service_interrupts = false;
```

```
      } else {
        out.state.can_service_interrupts = true;
      }
      out.state.lpre.set_have(true);
      out.instruction = in.instruction;
    }
  } else if (in.instruction.is_index()) {
    if (in.state.index.have) {
      out.state.can_service_interrupts = false;

      // Instruction was a NOP
      out.state.pre.have = false;
      out.state.lpre.have = false;
      out.state.index.have = false;
      out.instruction = nop;
    } else {
      if (in.state.pre.have || in.state.lpre.have) {
        out.state.can_service_interrupts = false;
      } else {
        out.state.can_service_interrupts = true;
      }
      out.state.index.have = true;
      out.instruction = in.instruction;
    }
  } else {
    if (
      in.state.pre.have || in.state.lpre.have || in.state.index.have
    ) {
      out.state.can_service_interrupts = false;
    } else {
      out.state.can_service_interrupts = true;
    }
    out.instruction = instruction;

    // Whenever we see an instruction other than pre, lpre, or index,
    // that means that those instructions stop being "in effect".
    out.state.pre.have = false;
    out.state.lpre.have = false;
    out.state.index.have = false;
  }
}
```

2.2. Instruction Group 1

The following encoding is used, with each character representing one bit:

`001i iiii oooo aaaa`, where

- `i` is a 5-bit sign-extended or zero-extended immediate, and is denoted `simm` when sign-extended or `imm` when zero-extended. Also, `simm` or `imm` can be expanded with `pre` or `lpre`.
- `a` encodes register `rA`
- `o` is the opcode

Here is a list of instructions from this encoding group.

- Opcode `0x0`: add rA, #simm
- Opcode `0x1`: add rA, pc, #simm
  - Effect: `rA <= pc + simm + 2;`
- Opcode `0x2`: add rA, sp, #simm
- Opcode `0x3`: add rA, fp, #simm
- Opcode `0x4`: cmp rA, #simm
  - Effect: Compare `rA` to `simm`.
  - Affectable flags: `Z`, `C`, `V`, `N`
- Opcode `0x5`: cpy rA, #simm
```

- Effect: Copy an immediate value into `rA`
- Opcode `0x6`: `lsl rA, #imm`
  - Effect: Logical shift left
- Opcode `0x7`: `lsr rA, #imm`
  - Effect: Logical shift right
- Opcode `0x8`: `asr rA, #imm`
  - Effect: Arithmetic shift right
- Opcode `0x9`: `and rA, #simm`
  - Effect: Bitwise AND
- Opcode `0xa`: `orr rA, #simm`
  - Effect: Bitwise OR
- Opcode `0xb`: `xor rA, #simm`
  - Effect: Bitwise XOR
- Opcode `0xc`: `ze rA, #imm`
  - Effect: Set `rA[31:imm]` to zero.
- Opcode `0xd`: `se rA, #imm`
  - Effect: Sign-extend `rA[imm:0]` to 32 bits, then copy that value to `rA`.
- Opcode `0xe`: `swi rA, #simm`
  - Effect: Call software interrupt number `rA + simm`.
- Opcode `0xf`: `swi #imm`
  - Effect: Call software interrupt number `imm`.

## 2.3. Instruction Group 2

The following encoding is used, with each character representing one bit:

`010f oooo bbbb aaaa`, where

- `o` is the opcode
- `b` encodes register `rB`
- `a` encodes register `rA`
- `f` is encoded as `0` if this instruction cannot affect flags and encoded `1` if this instruction is permitted to affect flags. Note that `cmp` is permitted to affect flags regardless of this bit.

Here is a list of instructions from this encoding group.
- Opcode `0x0`: `add rA, rB`
  - Mnemonic for when flags not affected: `add`
  - Mnemonic for when flags affected: `add.f`
  - Affectable flags: `Z`, `C`, `V`, `N`
- Opcode `0x1`: `sub rA, rB`
  - Mnemonic for when flags not affected: `sub`
  - Mnemonic for when flags affected: `sub.f`
  - Affectable flags: `Z`, `C`, `V`, `N`
- Opcode `0x2`: `add rA, sp, rB`
  - Mnemonic for when flags not affected: `add`
  - Mnemonic for when flags affected: `add.f`
  - Affectable flags: `Z`, `C`, `V`, `N`
- Opcode `0x3` `add rA, fp, rB`
  - Mnemonic for when flags not affected: `add`
  - Mnemonic for when flags affected: `add.f`
  - Affectable flags: `Z`, `C`, `V`, `N`
- Opcode `0x4`: `cmp rA, rB`
  - Effect: Compare `rA` to `rB`. `cmp` is *always* able to affect flags, independent of the encoded `f` bit of the instruction.
  - Affectable flags: `Z`, `C`, `V`, `N`
- Opcode `0x5`: `cpy rA, rB`

- Mnemonic for when flags not affected: `cpy`
- Mnemonic for when flags affected: `cpy.f`
- Effect: Copy `rB` into `rA`
- Affectable flags: `Z`, `N`
- Opcode `0x6`: `lsl rA, rB`
  - Mnemonic for when flags not affected: `lsl`
  - Mnemonic for when flags affected: `lsl.f`
  - Effect: Logical shift left
  - Affectable flags: `Z`, `N`
- Opcode `0x7`: `lsr rA, rB`
  - Mnemonic for when flags not affected: `lsr`
  - Mnemonic for when flags affected: `lsr.f`
  - Effect: Logical shift right
  - Affectable flags: `Z`, `N`
- Opcode `0x8`: `asr rA, rB`
  - Mnemonic for when flags not affected: `asr`
  - Mnemonic for when flags affected: `asr.f`
  - Effect: Arithmetic shift right
  - Affectable flags: `Z`, `N`
- Opcode `0x9`: `and rA, rB`
  - Mnemonic for when flags not affected: `and`
  - Mnemonic for when flags affected: `and.f`
  - Effect: Bitwise AND
  - Affectable flags: `Z`, `N`
- Opcode `0xa`: `orr rA, rB`
  - Mnemonic for when flags not affected: `orr`
  - Mnemonic for when flags affected: `orr.f`
  - Effect: Bitwise OR
  - Affectable flags: `Z`, `N`
- Opcode `0xb`: `xor rA, rB`
  - Mnemonic for when flags not affected: `xor`
  - Mnemonic for when flags affected: `xor.f`
  - Effect: Bitwise XOR
  - Affectable flags: `Z`, `N`
- Opcode `0xc`: `adc rA, rB`
  - Mnemonic for when flags not affected: `adc`
  - Mnemonic for when flags affected: `adc.f`
  - Effect: Add with Carry, using the formula `rA + rB + flags.C` to compute the value that will be written into `rA`.
  - Affectable flags: `Z`, `C`, `V`, `N`
- Opcode `0xd`: `sbc rA, rB`
  - Mnemonic for when flags not affected: `sbc`
  - Mnemonic for when flags affected: `sbc.f`
  - Effect: Subtract with Borrow, using the formula `rA + (~rB) + flags.C` to compute the value that will be written into `rA`.
  - Affectable flags: `Z`, `C`, `V`, `N`
- Opcode `0xe`: `cmpbc rA, rB`
  - Effect: Compare `rA` to `rB`, but with carry-in and a different effect for setting the `flags.Z`. `cmpbc` is *always* able to affect `flags`, independent of the encoded `f` bit of the instruction.
  - Note: this instruction acts much like `sbc rA, rB`, but without storing the subtraction's result into `rA`. However, this instruction sets the `Z` flag to `prev(flags.Z) AND ((rA + (~rB) + flags.C) == 0)`
  - Affectable flags: `Z`, `C`, `V`, `N`

## 2.4. Instruction Group 3: Relative Branches

The following encoding is used, with each character representing one bit:

`011i iiii iiii oooo`, where

- `i` is a 9-bit sign-extended immediate, which can be expanded by `pre` or `lpre`, and is denoted `simm`
- `o` is the opcode

Here is a list of instructions from this encoding group.
- Opcode `0x0`: `bl simm`
  - Name: Branch and Link
  - Description: Relative call
  - Effect: `lr <= pc + 2; pc <= pc + simm + 2;`
- Opcode `0x1`: `bra simm`
  - Name: BRanch Always
  - Description: Unconditional relative branch
  - Effect: `pc <= pc + simm + 2;`
- Opcode `0x2`: `beq simm`
  - Name: Branch if EQual
  - Effect: `if (flags.Z) pc <= pc + simm + 2;`
- Opcode `0x3`: `bne simm`
  - Name: Branch if Not Equal
  - Effect: `if (!flags.Z) pc <= pc + simm + 2;`
- Opcode `0x4`: `bmi simm`
  - Name: Branch if MInus
  - Effect: `if (flags.N) pc <= pc + simm + 2;`
- Opcode `0x5`: `bpl simm`
  - Name: Branch if PLus
  - Effect: `if (!flags.N) pc <= pc + simm + 2;`
- Opcode `0x6`: `bvs simm`
  - Name: Branch if oVerflow Set
  - Effect: `if (flags.V) pc <= pc + simm + 2;`
- Opcode `0x7`: `bvc simm`
  - Name: Branch if oVerflow Clear
  - Effect: `if (!flags.V) pc <= pc + simm + 2;`
- Opcode `0x8`: `bgeu simm`
  - Name: Branch if Greater than or Equal Unsigned
  - Effect: `if (flags.C) pc <= pc + simm + 2;`
- Opcode `0x9`: `bltu simm`
  - Name: Branch if Less Than Unsigned
  - Effect: `if (!flags.C) pc <= pc + simm + 2;`
- Opcode `0xa`: `bgtu simm`
  - Name: Branch if Greater Than Unsigned
  - Effect: `if (flags.C AND !flags.Z) pc <= pc + simm + 2;`
- Opcode `0xb`: `bleu simm`
  - Name: Branch if Less than or Equal Unsigned
  - Effect: `if (!flags.C OR flags.Z) pc <= pc + simm + 2;`
- Opcode `0xc`: `bges simm`
  - Name: Branch if Greater than or Equal Signed
  - Effect: `if (!(flags.N XOR flags.V)) pc <= pc + simm + 2;`
- Opcode `0xd`: `blts simm`
  - Name: Branch if Less Than Signed
  - Effect: `if (flags.N XOR flags.V) pc <= pc + simm + 2;`
- Opcode `0xe`: `bgts simm`
  - Name: Branch if Greater Than Signed

- Effect:

```
if (!(flags.N XOR flags.V) AND !flags.Z)
   pc <= pc + simm + 2;
```
- Opcode `0xf`: `bles simm`
  - Name: Branch if Less than or Equal Signed
  - Effect:

```
if ((flags.N XOR flags.V) OR flags.Z)
   pc <= pc + simm + 2;
```

2.5. Instruction Group 4

The following encoding is used, with each character representing one bit:

`100o oooo bbbb aaaa`, where

- `o` is the opcode
- `b` encodes register `rB` or register `sC`, where `sC` is one of the special-purpose registers
- `a` encodes register `rA` or register `sA`, where `sA` is one of the special-purpose registers

Here is a list of instructions from this encoding group.
- Opcode `0x0`: `jl rA`
  - Effect: `lr <= pc + 2; pc <= rA;`
- Opcode `0x1`: `jmp rA`
  - Effect: `pc <= rA;`

`jmp ira`
- Effect: `pc <= ira;`
- Opcode `0x3`: `reti`
  - Effect: enables IRQs (by copying `0x1` into `ie`) and performs `pc <= ira;`
- Opcode `0x4`: `ei`
  - Effect: copy `1` into `ie`.
- Opcode `0x5`: `di`
  - Effect: copy `0` into `ie`.
- Opcode `0x6`: `push rA, rB` `rB` as the stack pointer, post-decrementing `rB`.
  - This instruction does nothing when `rA` is the same register as `rB`.
  - Note: As a pseudo instruction, omitting ", `rB`" will automatically select `sp` as the particular stack pointer.
- Opcode `0x7`: `push sA, rB` `rB` as the stack pointer, post-decrementing `rB`.
  - Note that `sA` is considered to be 32-bit for the purpose of the store to memory and decrementing `rB`, even if `sA` is `flags` or `ie`.
  - Note: As a pseudo instruction, omitting ", `rB`" will automatically select `sp` as the particular stack pointer.
- Opcode `0x8`: `pop rA, rB` `rB` as the stack pointer, pre-incrementing `rB`.
  - This instruction does nothing when `rA` is the same register as `rB`.
  - Note: As a pseudo instruction, omitting ", `rB`" will automatically select `sp` as the particular stack pointer.
- Opcode `0x9`: `pop sA, rB`
  - Pseudo instruction with same effect:
    - `ldrib sA, rB`
    - Name explanation: load register, increment base before
  - Effect: pops `sA` off the stack, using `rB` as the stack pointer, pre-incrementing `rB`.
  - Note that `sA` is considered to be 32-bit for the purpose of the load from memory and incrementing `rB`, even if `sA` is `flags` or `ie`.
  - Note: As a pseudo instruction, omitting ", `rB`" will automatically select `sp` as the particular stack pointer.
- Opcode `0xa`: `pop pc, rB`

- Pseudo instruction with same effect:
  - `ldrib pc, rB`
  - Name explanation: load register, increment base before
- Effect: pops `pc` off the stack, using `rB` as the stack pointer, pre-incrementing `rB`.
- Note that `pc` is considered to be 32-bit for the purpose of the load from memory and incrementing `rB`, even if `pc` is `flags` or `ie`.
- Note: As a pseudo instruction, omitting ", `rB`" will automatically select `sp` as the particular stack pointer.
- Opcode `0xb`: `mul rA, rB`
  - Effect: `rA <= rA * rB;`
- Opcode `0xc`: `udiv rA, rB`
  - Effect: `rA <= u32(rA) / u32(rB);`
- Opcode `0xd`: `sdiv rA, rB`
  - Effect: `rA <= s32(rA) / s32(rB);`
- Opcode `0xe`: `umod rA, rB`
  - Effect: `rA <= u32(rA) \// u32(rB);`
- Opcode `0xf`: `smod rA, rB`
  - Effect: `rA <= s32(rA) \// s32(rB);`
- Opcode `0x10`: `lumul rA, rB`
  - Effect: This instruction multiplies `rA` by `rB`, performing an unsigned 32-bit by 32-bit -> 64-bit multiply, storing result in `concat{r0, r1}`.
- Opcode `0x11`: `lsmul rA, rB`
  - Effect: This instruction multiplies `rA` by `rB`, performing a signed 32-bit by 32-bit -> 64-bit multiply, storing result in `concat{r0, r1}`.
- Opcode `0x12`: `udiv64 rA, rB`
  - Effect: performs a 64-bit by 64-bit unsigned division of `concat{rA, r{A + 1}}` by `concat{rB, r{B + 1}}`, storing 64-bit result in `concat{rA, r{A + 1}}`.
  - Note: This instruction operates as if `rA` were encoded with (`A[0x0] == 0b0`) and `rB` were encoded with (`B[0x0] == 0b0`),
  - Note: This instruction executes more quickly if `rB`'s value is `0x00000000`, i.e. if the operation is actually a 64-bit by 32-bit -> 64-bit unsigned divide.
- Opcode `0x13`: `sdiv64 rA, rB`
  - Effect: performs a 64-bit by 64-bit signed division of `concat{rA, r{A + 1}}` by `concat{rB, r{B + 1}}`, storing 64-bit result in `concat{rA, r{A + 1}}`.
  - Note: This instruction operates as if `rA` were encoded with (`A[0x0] == 0b0`) and `rB` were encoded with (`B[0x0] == 0b0`),
  - Note: This instruction executes more quickly if `rB` is equal to bits `[63:32]` of `s64(r{B+1})`.
- Opcode `0x14`: `umod64 rA, rB`
  - Effect: performs a 64-bit by 64-bit unsigned modulo of `concat{rA, r{A + 1}}` by `concat{rB, r{B + 1}}`, storing 64-bit result in `concat{rA, r{A + 1}}`.
  - Note: This instruction operates as if `rA` were encoded with (`A[0x0] == 0b0`) and `rB` were encoded with (`B[0x0] == 0b0`),
  - Note: This instruction executes more quickly if `rB`'s value is `0x00000000`, i.e. if the operation is actually a 64-bit by 32-bit -> 64-bit unsigned modulo.
- Opcode `0x15`: `smod64 rA, rB`
  - Effect: performs a 64-bit by 64-bit signed modulo of `concat{rA, r{A + 1}}` by `concat{rB, r{B + 1}}`, storing 64-bit result in `concat{rA, r{A + 1}}`.
  - Note: This instruction operates as if `rA` were encoded with (`A[0x0] == 0b0`) and `rB` were encoded with (`B[0x0] == 0b0`),
  - Note: This instruction executes more quickly if `rB` is equal to bits `[63:32]` of `s64(r{B+1})`.
- Opcode `0x16`: `ldub rA, [rB]`
  - Effect: Load an 8-bit value from memory at address computed as `rB + index_reg`, zero-extend the 8-bit value to 32 bits, then put the zero-extended 32-bit value into `rA`.

- The `index_reg` value is guaranteed to be zero unless an `index` is in effect.
- Shorthand for having the assembler insert an `index rC` instruction before this one: `ldub rA, [rB, rC]`
- Opcode `0x17`: `ldsb rA, [rB]`
  - Effect: Load an 8-bit value from memory at address computed as `rB + index_reg`, sign-extend the 8-bit value to 32 bits, then put the sign-extended 32-bit value into `rA`.
  - The `index_reg` value is guaranteed to be zero unless an `index` is in effect.
  - Shorthand for having the assembler insert an `index rC` instruction before this one: `ldsb rA, [rB, rC]`
- Opcode `0x18`: `lduh rA, [rB]`
  - Effect: Load a 16-bit value from memory at address computed as `rB + index_reg`, zero-extend the 16-bit value to 32 bits, then put the zero-extended 32-bit value into `rA`.
  - The `index_reg` value is guaranteed to be zero unless an `index` is in effect.
  - Shorthand for having the assembler insert an `index rC` instruction before this one: `lduh rA, [rB, rC]`
- Opcode `0x19`: `ldsh rA, [rB]`
  - Effect: Load a 16-bit value from memory at address computed as `rB + index_reg`, sign-extend the 16-bit value to 32 bits, then put the zero-extended 32-bit value into `rA`.
  - The `index_reg` value is guaranteed to be zero unless an `index` is in effect.
  - Shorthand for having the assembler insert an `index rC` instruction before this one: `ldsh rA, [rB, rC]`
- Opcode `0x1a`: `stb rA, [rB]`
  - Effect: Store `rA[7:0]` to memory at the address computed as `rB + index_reg`.
  - The `index_reg` value is guaranteed to be zero unless an `index` is in effect.
  - Shorthand for having the assembler insert an `index rC` instruction before this one: `stb rA, [rB, rC]`
- Opcode `0x1b`: `sth rA, [rB]`
  - Effect: Store `rA[15:0]` to memory at the address computed as `rB + index_reg`.
  - The `index_reg` value is guaranteed to be zero unless an `index` is in effect.
  - Shorthand for having the assembler insert an `index rC` instruction before this one: `sth rA, [rB, rC]`
- Opcode `0x1c`: `cpy rA, sC`
  - Effect: `rA <= sC;`
- Opcode `0x1d`: `cpy sA, rB`
  - Effect: `sA <= rB;`
- Opcode `0x1e`: `cpy sA, sB`
  - Effect: `sA <= sB;`
- Opcode `0x1f`: `index rA`
  - Effect: Performs `index_reg <= rA;` and stores that `index` is in effect.
  - Note: If `index` is in effect and the current instruction is `index`, the current instruction will be treated as a NOP, and `index` will stop being in effect.
  - Note: `pre` and `index` can be combined with one another (though this is only useful for `ldr` and `str`).
  - Note: A non-`pre` instruction following `index` will store that that `index` is not in effect any more. (It will also store that `pre` is not in effect any more).
  - Note: If `index` is in effect, the current instruction cannot be interrupted by an IRQ.
  - Note: Any time `index` stops being in effect, `pre` will stop being in effect as well.

2.6. Instruction Group 5: Immediate Indexed Load

The following encoding is used, with each character representing one bit:

`101i iiii bbbb aaaa`, where

- `i` is a 5-bit sign-extended immediate, which can be expanded by `pre` or `lpre`, and is denoted `simm`
- `b` encodes register `rB`

- a encodes register `rA`

The one instruction from this encoding group is `ldr rA, [rB, #simm]`. This is a 32-bit load into `rA`, where the effective address to load from is computed as `rB + <index_reg> + simm`, using the sign-extended form of `simm`.

The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.

The following pseudo instructions exist within the assembler for this encoding group:
- `ldr rA, [rB]`
  - This pseudo instruction assembles to the following:
    - `ldr rA, [rB, #0x0]`
- `ldr rA, [rB, rC]`
  - This pseudo instruction assembles to the following:
    - `index rC`
    - `ldr rA, [rB, #0x0]`
- `ldr rA, [rB, rC, #simm]`
  - This pseudo instruction assembles to the following:
    - `index rC`
    - `ldr rA, [rB, #simm]`

2.7. Instruction Group 6: Immediate Indexed Store

The following encoding is used, with each character representing one bit:

`110i iiii bbbb aaaa`, where
- `i` is a 5-bit sign-extended immediate, which can be expanded by `pre` or `lpre`, and is denoted `simm`
- `b` encodes register `rB`
- `a` encodes register `rA`

The one instruction from this encoding group is `str rA, [rB, #simm]`. This is a 32-bit store of `rA`, where the effective address to store to is computed as `<index_reg> + rB + simm`, using the sign-extended form of `simm`.

The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.

The following pseudo instructions exist within the assembler for this encoding group:

- `str rA, [rB]`
  - This pseudo instruction assembles to the following:
    - `str rA, [rB, #0x0]`
- `str rA, [rB, rC]`
  - This pseudo instruction assembles to the following:
    - `index rC`
    - `str rA, [rB, #0x0]`
- `str rA, [rB, rC, #simm]`
  - This pseudo instruction assembles to the following:
    - `index rC`
    - `str rA, [rB, #simm]`

2.8. Instruction Group 7, Subgroup `0b00`:

Extra 8-bit and 16-bit Ops} The following encoding is used, with each character representing one bit:

`1110 0woo bbbb aaaa`, where
- `w` is the operation width
  - When `0b0`: 8-bit operation
  - When `0b0`: 16-bit operation
- `o` is the opcode
- `b` encodes register `rB`
- `a` encodes register `rA`

Here is a list of instructions from this encoding group.
- Opcode `0x0`:
  - w value `0b0`: `cmpb rA, rB`
    - Effect: Compare `rA[7:0]` to `rB[7:0]`
    - Affectable flags: `Z, C, V, N`
  - w value `0b1`: `cmph rA, rB`
    - Effect: Compare `rA[15:0]` to `rB[15:0]`
    - Affectable flags: `Z, C, V, N`
- Opcode `0x1`:
  - w value `0b0`: `lsrb rA, rB`
    - Effect: Logical shift right `rA[7:0]` by `rB`
  - w value `0b1`: `lsrh rA, rB`
    - Effect: Logical shift right `rA[15:0]` by `rB`
- Opcode `0x2`:
  - w value `0b0`: `asrb rA, rB`
    - Effect: Arithmetic shift right `rA[7:0]` by `rB`
  - w value `0b1`: `asrh rA, rB`
    - Effect: Arithmetic shift right `rA[15:0]` by `rB`

2.9. Instruction Group 7, Subgroup `0b010`: Extra load/store instructions

The following encoding is used, with each character representing one bit:
`1110 10oo bbbb aaaa`, where
- `o` is the opcode
- `b` encodes register `rB` or `sB`
- `a` encodes register `sA`

Here is a list of instructions from this encoding group.
- Opcode `0x0`: `ldr sA, [rB]`
  - Effect: 32-bit load of `sA` from memory at address held in `rB`.
- Opcode `0x1`: `ldr sA, [sB]`
  - Effect: 32-bit load of `sA` from memory at address held in `sB`.
- Opcode `0x2`: `str sA, [rB]`
  - Effect: 32-bit store of `sA` to memory at address held in `rB`.
- Opcode `0x3`: `str sA, [sB]`
  - Effect: 32-bit store of `sA` to memory at address held in `sB`.

2.10. Instruction Group 7, Subgroup `0b0110`: Icache Reload Instruction

The following encoding is used, with each character representing one bit:
`1110 110i iiii aaaa`, where
- `i` is a 5-bit sign-extended immediate, which can be expanded by `pre` or `lpre`, and is denoted `simm`.
- `a` encodes register `rA`

The one instruction from this encoding group is `icreload [rA, #simm]`. This instruction forcibly reloads the icache line at an effective address computed as `<index_reg> + rA + simm`, using the sign-extended form of `simm`. As a side note, this instruction will attempt to read the data from dcache, and will only read from memory if dcache doesn't have that data. Icache lines and dcache lines are the same length, 64 bytes.

The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.

The following pseudo instructions exist within the assembler for this encoding group:
- `icreload [rA]`
  - This pseudo instruction assembles to the following:
    - `icreload [rA, #0x0]`
- `icreload [rA, rC]`

- This pseudo instruction assembles to the following:
  - `index rC`
  - `icreload [rA, #0x0]`
- `icreload [rA, rC, #simm]`
  - This pseudo instruction assembles to the following:
    - `index rC`
    - `icreload [rA, #simm]`