

Flare32 CPU

FL4SHK

May 26, 2023

Table of Contents

Table of Contents	1
Introduction	2
2.1 General Information	2
2.2 Registers	2
Instruction Set	4
3.1 Instruction Group 0: pre and lpre	4
3.1.1 Handling of pre , lpre , and index	4
3.2 Instruction Group 1	6
3.3 Instruction Group 2	8
3.4 Instruction Group 3: Relative Branches	11
3.5 Instruction Group 4	13
3.6 Instruction Group 5: Immediate Indexed Load	18
3.7 Instruction Group 6: Immediate Indexed Store	19
3.8 Instruction Group 7, Subgroup 0b00: Extra 8-bit and 16-bit Ops	20
3.9 Instruction Group 7, Subgroup 0b010: Extra load/store instructions	21
3.10 Instruction Group 7, Subgroup 0b0110: Icache Reload Instruction	22

Introduction

2.1 General Information

- Addresses are 32-bit.
- Big-endian byte ordering is used.
- Bytes are octets (8 bits).
- Instructions must be aligned to 16 bits, so jump and branch targets must also be aligned to 16 bits.
 - Branch offsets encoded into instructions must be 16-bit aligned, or in other words bit 0 of the branch offset must be 0b0

2.2 Registers

There are sixteen general-purpose registers (all of which are 32-bit): `r0`, `r1`, `r2`, ..., `r11`, `r12`, `lr` (link register, the return address of `bl` and `jl`), `fp` (frame pointer), `sp` (stack pointer).

The program counter, `pc`, is 32 bits long, as addresses are 32-bit.
Here are the special-purpose registers:

- `flags`: arithmetic/logic FLAGS; (reg encoding: `0x0`)
- `ids`: Interrupt DeStination: the address to jump to upon an interrupt being serviced (also known as the interrupt vector); (reg encoding: `0x1`)
- `ira`: Interrupt Return Address: the address that was jumped from upon an interrupt being serviced; (reg encoding: `0x2`)
- `ie`: IRQ Enable flag: flag indicating whether IRQs are disabled (`0x0`) or enabled (`0x1`); note that this flag starts with a value of `0x0`; (reg encoding: `0x3`)
- `ity`: Interrupt TYpe: flag indicating whether the most recently taken interrupt is an IRQ (`0x0`) or a `swi` (`0x1`) (reg encoding: `0x4`)
- `sty`: Software interrupt TYpe: `swi`'s argument. For `swi rA, #simm`, this is the value `rA + simm` (reg encoding: `0x5`)
- Note: All other encodings for special-purpose registers are reserved.

Here are the bits of `flags`:

- Zero (Z): (`flags` bit 0)

- Carry (C): (**flags** bit 1)
- oVerflow (V): (**flags** bit 2)
- Negative (N): (**flags** bit 3)
- Note: All other bit positions of **flags** are reserved.

Instruction Set

3.1 Instruction Group 0: **pre** and **lpre**

For **pre**, the following encoding is used, with each character representing one bit:

0000 *iiii* *iiii* *iiii*, where

- *i* is a 12-bit constant.

For **lpre**, the following encoding is used, with each character representing one bit:

0001 0*iiii* *iiii* *iiii* *iiii* *iiii* *iiii* *iiii*, where

- *i* is a 27-bit constant.

pre and **lpre** are mechanisms by which immediates larger than normal can be used, essentially acting like variable width instructions.

There is no mechanism in the assembly language itself to use **pre** or **lpre** as instructions. Instead, it is expected that the assembler or linker will be the one to insert **pre** or **lpre** as needed if an immediate is too large for a particular instruction.

For non-branch instructions:

- If **pre** is used, the immediate field of the **pre** instruction will form bits [16:5] of the immediate of the next non-**index** instruction. The 17-bit immediate will then be sign-extended to 32 bits. Sign-extension will be performed on the 17-bit immediate even if, had there been no **pre**, the 5-bit immediate would have been zero-extended.
- On the other hand, if **lpre** is used, the immediate field of the **lpre** instruction will form bits [31:5] of the immediate of the next non-**index** instruction.

For branch instructions (group 3):

- if **pre** is used, the immediate field of the **pre** instruction will form bits [20:9] of the immediate of the next non-**index** instruction. The 21-bit immediate will then be sign-extended to 32 bits.
- On the other hand, if **lpre** is used, the immediate field of the **lpre** instruction will form bits [31:9] of the immediate of the next non-**index** instruction.

3.1.1 Handling of **pre**, **lpre**, and **index**

When a **pre** or **lpre** instruction is found, **pre** or **lpre** will be considered to be "in effect". This condition lasts for one or two instructions after the **pre** or **lpre** instruction, depending on whether or not **index** was in effect.

`index` is an instruction (defined later) that allows a subsequent load or store instruction to perform `base_reg + index_reg` indexing. When an `index` instruction is found, it is considered to be in effect. Like `pre` and `lpre`, `index` is inserted automatically by the assembler.

`index` can be combined with `pre` or `lpre`, and it does not matter whether `index` or the `pre/lpre` instruction came first.

When `pre`, `lpre`, or `index` is in effect, IRQs will not be serviced.

Pseudo code for handling the how `pre`, `lpre`, and `index` are to be handled regarding whether or not they are "in effect" is as follows:

```
function handle_pre_lpre_index(input in, output out) {
    if (in.instruction.is_pre()) {
        if (in.state.pre.have || in.state.lpre.have) {
            out.state.can_service_interrupts = false;

            // Instruction was a NOP
            out.state.pre.have = false;
            out.state.lpre.have = false;
            out.state.index.have = false;
            out.instruction = nop;
        } else {
            if (in.state.index.have) {
                out.state.can_service_interrupts = false;
            } else {
                out.state.can_service_interrupts = true;
            }
            out.state.pre.set_have(true);
            out.instruction = in.instruction;
        }
    } else if (in.instruction.is_lpre()) {
        if (in.state.pre.have || in.state.lpre.have) {
            out.state.can_service_interrupts = false;

            // Instruction was a NOP
            out.state.pre.have = false;
            out.state.lpre.have = false;
            out.state.index.have = false;
            out.instruction = nop;
        } else {
            if (in.state.index.have) {
                out.state.can_service_interrupts = false;
            } else {
                out.state.can_service_interrupts = true;
            }
            out.state.lpre.set_have(true);
            out.instruction = in.instruction;
        }
    }
}
```

```

    }
} else if (in.instruction.is_index()) {
    if (in.state.index.have) {
        out.state.can_service_interrupts = false;

        // Instruction was a NOP
        out.state.pre.have = false;
        out.state.lpre.have = false;
        out.state.index.have = false;
        out.instruction = nop;
    } else {
        if (in.state.pre.have || in.state.lpre.have) {
            out.state.can_service_interrupts = false;
        } else {
            out.state.can_service_interrupts = true;
        }
        out.state.index.have = true;
        out.instruction = in.instruction;
    }
} else {
    if (
        in.state.pre.have || in.state.lpre.have || in.state.index.have
    ) {
        out.state.can_service_interrupts = false;
    } else {
        out.state.can_service_interrupts = true;
    }
    out.instruction = instruction;

    // Whenever we see an instruction other than pre, lpre, or index,
    // that means that those instructions stop being "in effect".
    out.state.pre.have = false;
    out.state.lpre.have = false;
    out.state.index.have = false;
}
}

```

3.2 Instruction Group 1

The following encoding is used, with each character representing one bit:
001i iiii oooo aaaa, where

- i is a 5-bit sign-extended or zero-extended immediate, and is denoted **simm** when sign-extended or **imm** when zero-extended. Also, **simm** or **imm** can be expanded with **pre** or **lpre**.

- `a` encodes register `rA`
- `o` is the opcode

Here is a list of instructions from this encoding group.

- Opcode 0x0: `add rA, #simm`
- Opcode 0x1: `add rA, pc, #simm`
 - Effect: `rA <= pc + simm + 2;`
- Opcode 0x2: `add rA, sp, #simm`
- Opcode 0x3: `add rA, fp, #simm`
- Opcode 0x4: `cmp rA, #simm`
 - Effect: Compare `rA` to `simm`.
 - Affectable flags: Z, C, V, N
- Opcode 0x5: `cpy rA, #simm`
 - Effect: Copy an immediate value into `rA`
- Opcode 0x6: `lsl rA, #imm`
 - Effect: Logical shift left
- Opcode 0x7: `lsr rA, #imm`
 - Effect: Logical shift right
- Opcode 0x8: `asr rA, #imm`
 - Effect: Arithmetic shift right
- Opcode 0x9: `and rA, #simm`
 - Effect: Bitwise AND
- Opcode 0xa: `orr rA, #simm`
 - Effect: Bitwise OR
- Opcode 0xb: `xor rA, #simm`
 - Effect: Bitwise XOR
- Opcode 0xc: `ze rA, #imm`
 - Effect: Set `rA[31:simm]` to zero.

- Opcode 0xd: **se** *rA*, #imm
 - Effect: Sign-extend *rA*[simm:0] to 32 bits, then copy that value to *rA*.
- Opcode 0xe: **swi** *rA*, #simm
 - Effect: Call software interrupt number *rA* + *simm*.
- Opcode 0xf: **swi** #imm
 - Effect: Call software interrupt number *imm*.

3.3 Instruction Group 2

The following encoding is used, with each character representing one bit:

010f oooo bbbb aaaa, where

- o is the opcode
- b encodes register *rB*
- a encodes register *rA*
- f is encoded as 0 if this instruction cannot affect flags and encoded 1 if this instruction is permitted to affect flags. Note that **cmp** is permitted to affect flags regardless of this bit.

Here is a list of instructions from this encoding group.

- Opcode 0x0: **add** *rA*, *rB*
 - Mnemonic for when flags not affected: **add**
 - Mnemonic for when flags affected: **add.f**
 - Affectable flags: Z, C, V, N
- Opcode 0x1: **sub** *rA*, *rB*
 - Mnemonic for when flags not affected: **sub**
 - Mnemonic for when flags affected: **sub.f**
 - Affectable flags: Z, C, V, N
- Opcode 0x2: **add** *rA*, *sp*, *rB*
 - Mnemonic for when flags not affected: **add**
 - Mnemonic for when flags affected: **add.f**

- Affectable flags: Z, C, V, N
- Opcode 0x3: **add rA, fp, rB**
 - Mnemonic for when flags not affected: **add**
 - Mnemonic for when flags affected: **add.f**
 - Affectable flags: Z, C, V, N
- Opcode 0x4: **cmp rA, rB**
 - Effect: Compare **rA** to **rB**. **cmp** is *always* able to affect flags, independent of the encoded **f** bit of the instruction.
 - Affectable flags: Z, C, V, N
- Opcode 0x5: **cpy rA, rB**
 - Mnemonic for when flags not affected: **cpy**
 - Mnemonic for when flags affected: **cpy.f**
 - Effect: Copy **rB** into **rA**
 - Affectable flags: Z, N
- Opcode 0x6: **lsl rA, rB**
 - Mnemonic for when flags not affected: **lsl**
 - Mnemonic for when flags affected: **lsl.f**
 - Effect: Logical shift left
 - Affectable flags: Z, N
- Opcode 0x7: **lsr rA, rB**
 - Mnemonic for when flags not affected: **lsr**
 - Mnemonic for when flags affected: **lsr.f**
 - Effect: Logical shift right
 - Affectable flags: Z, N
- Opcode 0x8: **asr rA, rB**
 - Mnemonic for when flags not affected: **asr**
 - Mnemonic for when flags affected: **asr.f**
 - Effect: Arithmetic shift right
 - Affectable flags: Z, N
- Opcode 0x9: **and rA, rB**
 - Mnemonic for when flags not affected: **and**
 - Mnemonic for when flags affected: **and.f**

- Effect: Bitwise AND
- Affectable flags: Z, N
- Opcode 0xa: **orr rA, rB**
 - Mnemonic for when flags not affected: **orr**
 - Mnemonic for when flags affected: **orr.f**
 - Effect: Bitwise OR
 - Affectable flags: Z, N
- Opcode 0xb: **xor rA, rB**
 - Mnemonic for when flags not affected: **xor**
 - Mnemonic for when flags affected: **xor.f**
 - Effect: Bitwise XOR
 - Affectable flags: Z, N
- Opcode 0xc: **adc rA, rB**
 - Mnemonic for when flags not affected: **adc**
 - Mnemonic for when flags affected: **adc.f**
 - Effect: Add with Carry, using the formula $rA + rB + \text{flags.C}$ to compute the value that will be written into rA.
 - Affectable flags: Z, C, V, N
- Opcode 0xd: **sbc rA, rB**
 - Mnemonic for when flags not affected: **sbc**
 - Mnemonic for when flags affected: **sbc.f**
 - Effect: Subtract with Borrow, using the formula $rA - (rB) + \text{flags.C}$ to compute the value that will be written into rA.
 - Affectable flags: Z, C, V, N
- Opcode 0xe: **cmpbc rA, rB**
 - Effect: Compare rA to rB, but with carry-in and a different effect for setting the **flags.Z**. **cmpbc** is *always* able to affect **flags**, independent of the encoded **f** bit of the instruction.
 - Note: this instruction acts much like **sbc rA, rB**, but without storing the subtraction's result into rA. However, this instruction sets the Z flag to $\text{prev}(\text{flags.Z}) \text{ AND } ((rA - (rB) + \text{flags.C}) == 0)$
 - Affectable flags: Z, C, V, N

3.4 Instruction Group 3: Relative Branches

The following encoding is used, with each character representing one bit:

011i iiii iiii oooo, where

- i is a 9-bit sign-extended immediate, which can be expanded by `pre` or `lpre`, and is denoted `simm`
- o is the opcode

Here is a list of instructions from this encoding group.

- Opcode 0x0: **bl** `simm`
 - Name: Branch and Link
 - Description: Relative call
 - Effect: `lr <= pc + 2; pc <= pc + simm + 2;`
- Opcode 0x1: **bra** `simm`
 - Name: BRanch Always
 - Description: Unconditional relative branch
 - Effect: `pc <= pc + simm + 2;`
- Opcode 0x2: **beq** `simm`
 - Name: Branch if EQual
 - Effect: `if (flags.Z) pc <= pc + simm + 2;`
- Opcode 0x3: **bne** `simm`
 - Name: Branch if Not Equal
 - Effect: `if (!flags.Z) pc <= pc + simm + 2;`
- Opcode 0x4: **bmi** `simm`
 - Name: Branch if MInus
 - Effect: `if (flags.N) pc <= pc + simm + 2;`
- Opcode 0x5: **bpl** `simm`
 - Name: Branch if PLus
 - Effect: `if (!flags.N) pc <= pc + simm + 2;`
- Opcode 0x6: **bvs** `simm`
 - Name: Branch if oVerflow Set

- Effect: if (flags.V) $pc \leq pc + \text{sim}m + 2$;
- Opcode 0x7: **bvc** *sim*m
 - Name: Branch if oVerflow Clear
 - Effect: if (!flags.V) $pc \leq pc + \text{sim}m + 2$;
- Opcode 0x8: **bgeu** *sim*m
 - Name: Branch if Greater than or Equal Unsigned
 - Effect: if (flags.C) $pc \leq pc + \text{sim}m + 2$;
- Opcode 0x9: **bltu** *sim*m
 - Name: Branch if Less Than Unsigned
 - Effect: if (!flags.C) $pc \leq pc + \text{sim}m + 2$;
- Opcode 0xa: **bgtu** *sim*m
 - Name: Branch if Greater Than Unsigned
 - Effect: if (flags.C AND !flags.Z) $pc \leq pc + \text{sim}m + 2$;
- Opcode 0xb: **bleu** *sim*m
 - Name: Branch if Less than or Equal Unsigned
 - Effect: if (!flags.C OR flags.Z) $pc \leq pc + \text{sim}m + 2$;
- Opcode 0xc: **bges** *sim*m
 - Name: Branch if Greater than or Equal Signed
 - Effect: if (!(flags.N XOR flags.V)) $pc \leq pc + \text{sim}m + 2$;
- Opcode 0xd: **blts** *sim*m
 - Name: Branch if Less Than Signed
 - Effect: if (flags.N XOR flags.V) $pc \leq pc + \text{sim}m + 2$;
- Opcode 0xe: **bgts** *sim*m
 - Name: Branch if Greater Than Signed
 - Effect: if (!(flags.N XOR flags.V) AND !flags.Z) $pc \leq pc + \text{sim}m + 2$;
- Opcode 0xf: **bles** *sim*m
 - Name: Branch if Less than or Equal Signed
 - Effect: if ((flags.N XOR flags.V) OR flags.Z) $pc \leq pc + \text{sim}m + 2$;

3.5 Instruction Group 4

The following encoding is used, with each character representing one bit:

100o oooo bbbb aaaa, where

- o is the opcode
- b encodes register **rB** or register **sC**, where **sC** is one of the special-purpose registers
- a encodes register **rA** or register **sA**, where **sA** is one of the special-purpose registers

Here is a list of instructions from this encoding group.

- Opcode 0x0: **j1 rA**
 - Effect: $lr \leq pc + 2$; $pc \leq rA$;
- Opcode 0x1: **jmp rA**
 - Effect: $pc \leq rA$;
- Opcode 0x2: **jmp ira**
 - Effect: $pc \leq ira$;
- Opcode 0x3: **reti**
 - Effect: enables IRQs (by copying 0x1 into **ie**) and performs $pc \leq ira$;
- Opcode 0x4: **ei**
 - Effect: copy 1 into **ie**.
- Opcode 0x5: **di**
 - Effect: copy 0 into **ie**.
- Opcode 0x6: **push rA, rB**
 - Effect: pushes **rA** onto the stack, using **rB** as the stack pointer, post-decrementing **rB**.
 - This instruction does nothing when **rA** is the same register as **rB**.
 - Note: As a pseudo instruction, omitting ", **rB**" will automatically select **sp** as the particular stack pointer. **sp**
- Opcode 0x7: **push sA, rB**

- Effect: pushes **sA** onto the stack, using **rB** as the stack pointer, post-decrementing **rB**.
 - Note that **sA** is considered to be 32-bit for the purpose of the store to memory and decrementing **rB**, even if **sA** is **flags** or **ie**.
 - Note: As a pseudo instruction, omitting ", **rB**" will automatically select **sp** as the particular stack pointer.
- Opcode 0x8: **pop rA, rB**
 - Effect: pops **rA** off the stack, using **rB** as the stack pointer, pre-incrementing **rB**.
 - This instruction does nothing when **rA** is the same register as **rB**.
 - Note: As a pseudo instruction, omitting ", **rB**" will automatically select **sp** as the particular stack pointer.
- Opcode 0x9: **pop sA, rB**
 - Pseudo instruction with same effect:
 - * **ldrib sA, rB**
 - * Name explanation: load register, increment base before
 - Effect: pops **sA** off the stack, using **rB** as the stack pointer, pre-incrementing **rB**.
 - Note that **sA** is considered to be 32-bit for the purpose of the load from memory and incrementing **rB**, even if **sA** is **flags** or **ie**.
 - Note: As a pseudo instruction, omitting ", **rB**" will automatically select **sp** as the particular stack pointer.
- Opcode 0xa: **pop pc, rB**
 - Pseudo instruction with same effect:
 - * **ldrib pc, rB**
 - * Name explanation: load register, increment base before
 - Effect: pops **pc** off the stack, using **rB** as the stack pointer, pre-incrementing **rB**.
 - Note that **pc** is considered to be 32-bit for the purpose of the load from memory and incrementing **rB**, even if **pc** is **flags** or **ie**.
 - Note: As a pseudo instruction, omitting ", **rB**" will automatically select **sp** as the particular stack pointer.
- Opcode 0xb: **mul rA, rB**
 - Effect: **rA <= rA * rB;**
- Opcode 0xc: **udiv rA, rB**

- Effect: $rA \leftarrow u32(rA) / u32(rB);$
- Opcode 0xd: **sdiv** *rA*, *rB*
 - Effect: $rA \leftarrow s32(rA) / s32(rB);$
- Opcode 0xe: **umod** *rA*, *rB*
 - Effect: $rA \leftarrow u32(rA) \% u32(rB);$
- Opcode 0xf: **smod** *rA*, *rB*
 - Effect: $rA \leftarrow s32(rA) \% s32(rB);$
- Opcode 0x10: **lumul** *rA*, *rB*
 - Effect: This instruction multiplies *rA* by *rB*, performing an unsigned 32-bit by 32-bit -> 64-bit multiply, storing result in `concat{r0, r1}`.
- Opcode 0x11: **lsmul** *rA*, *rB*
 - Effect: This instruction multiplies *rA* by *rB*, performing a signed 32-bit by 32-bit -> 64-bit multiply, storing result in `concat{r0, r1}`.
- Opcode 0x12: **udiv64** *rA*, *rB*
 - Effect: performs a 64-bit by 64-bit unsigned division of `concat{rA, r{A + 1}}` by `concat{rB, r{B + 1}}`, storing 64-bit result in `concat{rA, r{A + 1}}`.
 - Note: This instruction operates as if *rA* were encoded with (`A[0x0] == 0b0`) and *rB* were encoded with (`B[0x0] == 0b0`),
 - Note: This instruction executes more quickly if *rB*'s value is 0x00000000, i.e. if the operation is actually a 64-bit by 32-bit -> 64-bit unsigned divide.
- Opcode 0x13: **sdiv64** *rA*, *rB*
 - Effect: performs a 64-bit by 64-bit signed division of `concat{rA, r{A + 1}}` by `concat{rB, r{B + 1}}`, storing 64-bit result in `concat{rA, r{A + 1}}`.
 - Note: This instruction operates as if *rA* were encoded with (`A[0x0] == 0b0`) and *rB* were encoded with (`B[0x0] == 0b0`),
 - Note: This instruction executes more quickly if *rB* is equal to bits [63:32] of `s64(r{B+1})`.
- Opcode 0x14: **umod64** *rA*, *rB*
 - Effect: performs a 64-bit by 64-bit unsigned modulo of `concat{rA, r{A + 1}}` by `concat{rB, r{B + 1}}`, storing 64-bit result in `concat{rA, r{A + 1}}`.

- Note: This instruction operates as if **rA** were encoded with (**A[0x0]** == **0b0**) and **rB** were encoded with (**B[0x0]** == **0b0**),
- Note: This instruction executes more quickly if **rB**'s value is **0x00000000**, i.e. if the operation is actually a 64-bit by 32-bit -> 64-bit unsigned modulo.
- Opcode **0x15: smod64 rA, rB**
 - Effect: performs a 64-bit by 64-bit signed modulo of **concat{rA, r{A + 1}}** by **concat{rB, r{B + 1}}**, storing 64-bit result in **concat{rA, r{A + 1}}**.
 - Note: This instruction operates as if **rA** were encoded with (**A[0x0]** == **0b0**) and **rB** were encoded with (**B[0x0]** == **0b0**),
 - Note: This instruction executes more quickly if **rB** is equal to bits **[63:32]** of **s64(r{B+1})**.
- Opcode **0x16: ldub rA, [rB]**
 - Effect: Load an 8-bit value from memory at address computed as **rB + <index_reg>**, zero-extend the 8-bit value to 32 bits, then put the zero-extended 32-bit value into **rA**.
 - The **<index_reg>** value is guaranteed to be zero unless an **index** is in effect.
 - Shorthand for having the assembler insert an **index rC** instruction before this one: **ldub rA, [rB, rC]**
- Opcode **0x17: ldsb rA, [rB]**
 - Effect: Load an 8-bit value from memory at address computed as **rB + <index_reg>**, sign-extend the 8-bit value to 32 bits, then put the sign-extended 32-bit value into **rA**.
 - The **<index_reg>** value is guaranteed to be zero unless an **index** is in effect.
 - Shorthand for having the assembler insert an **index rC** instruction before this one: **ldsb rA, [rB, rC]**
- Opcode **0x18: ldubh rA, [rB]**
 - Effect: Load a 16-bit value from memory at address computed as **rB + <index_reg>**, zero-extend the 16-bit value to 32 bits, then put the zero-extended 32-bit value into **rA**.
 - The **<index_reg>** value is guaranteed to be zero unless an **index** is in effect.
 - Shorthand for having the assembler insert an **index rC** instruction before this one: **ldubh rA, [rB, rC]**

- Opcode 0x19: **ldsh** rA, [rB]
 - Effect: Load a 16-bit value from memory at address computed as rB + <index_reg>, sign-extend the 16-bit value to 32 bits, then put the zero-extended 32-bit value into rA.
 - The <index_reg> value is guaranteed to be zero unless an **index** is in effect.
 - Shorthand for having the assembler insert an **index** rC instruction before this one: **ldsh** rA, [rB, rC]
- Opcode 0x1a: **stb** rA, [rB]
 - Effect: Store rA[7:0] to memory at the address computed as rB + <index_reg>.
 - The <index_reg> value is guaranteed to be zero unless an **index** is in effect.
 - Shorthand for having the assembler insert an **index** rC instruction before this one: **stb** rA, [rB, rC]
- Opcode 0x1b: **sth** rA, [rB]
 - Effect: Store rA[15:0] to memory at the address computed as rB + <index_reg>.
 - The <index_reg> value is guaranteed to be zero unless an **index** is in effect.
 - Shorthand for having the assembler insert an **index** rC instruction before this one: **sth** rA, [rB, rC]
- Opcode 0x1c: **cpy** rA, sC
 - Effect: rA <= sC;
- Opcode 0x1d: **cpy** sA, rB
 - Effect: sA <= rB;
- Opcode 0x1e: **cpy** sA, sB
 - Effect: sA <= sB;
- Opcode 0x1f: **index** rA
 - Effect: Performs <index_reg> <= rA; and stores that **index** is in effect.
 - Note: If **index** is in effect and the current instruction is **index**, the current instruction will be treated as a NOP, and **index** will stop being in effect.

- Note: **pre** and **index** can be combined with one another (though this is only useful for **ldr** and **str**).
- Note: A non-**pre** instruction following **index** will store that that **index** is not in effect any more. (It will also store that **pre** is not in effect any more).
- Note: If **index** is in effect, the current instruction cannot be interrupted by an IRQ.
- Note: Any time **index** stops being in effect, **pre** will stop being in effect as well.

3.6 Instruction Group 5: Immediate Indexed Load

The following encoding is used, with each character representing one bit:

101i iiii bbbb aaaa, where

- **i** is a 5-bit sign-extended immediate, which can be expanded by **pre** or **lpre**, and is denoted **simm**
- **b** encodes register **rB**
- **a** encodes register **rA**

The one instruction from this encoding group is **ldr rA, [rB, #simm]**. This is a 32-bit load into **rA**, where the effective address to load from is computed as **rB + <index_reg> + simm**, using the sign-extended form of **simm**.

The **<index_reg>** value is guaranteed to be zero unless an **index** is in effect.

The following pseudo instructions exist within the assembler for this encoding group:

- **ldr rA, [rB]**
 - This pseudo instruction assembles to the following:


```
* ldr rA, [rB, #0x0]
```
- **ldr rA, [rB, rC]**
 - This pseudo instruction assembles to the following:

- * **index** rC
- * **ldr** rA, [rB, #0x0]
- **ldr** rA, [rB, rC, #simm]

– This pseudo instruction assembles to the following:

- * **index** rC
- * **ldr** rA, [rB, #simm]

3.7 Instruction Group 6: Immediate Indexed Store

The following encoding is used, with each character representing one bit:

110i iiii bbbb aaaa, where

- i is a 5-bit sign-extended immediate, which can be expanded by **pre** or **lpre**, and is denoted **simm**
- b encodes register rB
- a encodes register rA

The one instruction from this encoding group is **str** rA, [rB, #simm]. This is a 32-bit store of rA, where the effective address to store to is computed as $\langle \text{index_reg} \rangle + \text{rB} + \text{simm}$, using the sign-extended form of **simm**.

The $\langle \text{index_reg} \rangle$ value is guaranteed to be zero unless an **index** is in effect.

The following pseudo instructions exist within the assembler for this encoding group:

- **str** rA, [rB]
- This pseudo instruction assembles to the following:
 - * **str** rA, [rB, #0x0]
- **str** rA, [rB, rC]

– This pseudo instruction assembles to the following:

- * **index** rC
- * **str** rA, [rB, #0x0]
- **str** rA, [rB, rC, #simm]
 - This pseudo instruction assembles to the following:
 - * **index** rC
 - * **str** rA, [rB, #simm]

3.8 Instruction Group 7, Subgroup 0b00: Extra 8-bit and 16-bit Ops

The following encoding is used, with each character representing one bit:

1110 0w00 bbbb aaaa, where

- w is the operation width
 - When 0b0: 8-bit operation
 - When 0b1: 16-bit operation
- o is the opcode
- b encodes register rB
- a encodes register rA

Here is a list of instructions from this encoding group.

- Opcode 0x0:
 - w value 0b0: **cmpb** rA, rB
 - * Effect: Compare rA[7:0] to rB[7:0]
 - * Affectable flags: Z, C, V, N
 - w value 0b1: **cmph** rA, rB
 - * Effect: Compare rA[15:0] to rB[15:0]
 - * Affectable flags: Z, C, V, N
- Opcode 0x1:
 - w value 0b0: **lsrb** rA, rB

- * Effect: Logical shift right `rA[7:0]` by `rB`
- `w` value 0b1: `lsrh rA, rB`
 - * Effect: Logical shift right `rA[15:0]` by `rB`
- Opcode 0x2:
 - `w` value 0b0: `asrb rA, rB`
 - * Effect: Arithmetic shift right `rA[7:0]` by `rB`
 - `w` value 0b1: `asrh rA, rB`
 - * Effect: Arithmetic shift right `rA[15:0]` by `rB`

3.9 Instruction Group 7, Subgroup 0b010: Extra load/store instructions

The following encoding is used, with each character representing one bit:

1110 10oo bbbb aaaa, where

- `o` is the opcode
- `b` encodes register `rB` or `sB`
- `a` encodes register `sA`

Here is a list of instructions from this encoding group.

- Opcode 0x0: `ldr sA, [rB]`
 - Effect: 32-bit load of `sA` from memory at address held in `rB`.
- Opcode 0x1: `ldr sA, [sB]`
 - Effect: 32-bit load of `sA` from memory at address held in `sB`.
- Opcode 0x2: `str sA, [rB]`
 - Effect: 32-bit store of `sA` to memory at address held in `rB`.
- Opcode 0x3: `str sA, [sB]`
 - Effect: 32-bit store of `sA` to memory at address held in `sB`.

3.10 Instruction Group 7, Subgroup 0b0110: Icache Reload Instruction

The following encoding is used, with each character representing one bit:
1110 110i iiii aaaa, where

- *i* is a 5-bit sign-extended immediate, which can be expanded by *pre* or *lpre*, and is denoted *simm*.
- *a* encodes register *rA*

The one instruction from this encoding group is **icreload** [*rA*, #*simm*]. This instruction forcibly reloads the icache line at an effective address computed as `<index_reg> + rA + simm`, using the sign-extended form of *simm*. As a side note, this instruction will attempt to read the data from dcache, and will only read from memory if dcache doesn't have that data. Icache lines and dcache lines are the same length, 64 bytes.

The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.

The following pseudo instructions exist within the assembler for this encoding group:

- **icreload** [*rA*]
 - This pseudo instruction assembles to the following:

```
* icreload [rA, #0x0]
```
- **icreload** [*rA*, *rC*]
 - This pseudo instruction assembles to the following:

```
* index rC  
* icreload [rA, #0x0]
```
- **icreload** [*rA*, *rC*, #*simm*]
 - This pseudo instruction assembles to the following:

```
* index rC  
* icreload [rA, #simm]
```