

I Introduction

I.A General Information

- Addresses are 32-bit.
- Little-endian byte ordering is used.
- Bytes are octets (8 bits).
- Instructions must be aligned to 16 bits, so jump and branch targets must also be aligned to 16 bits.
 - Branch offsets encoded into instructions must be 16-bit aligned, or in other words bit 0 of the branch offset must be `0b0`

I.B Registers

There are sixteen general-purpose registers (all of which are 32-bit): `r0`, `r1`, `r2`, ..., `r11`, `r12`, `lr` (link register, the return address of `bl` and `jl`), `fp` (frame pointer), and `sp` (stack pointer).

The program counter, `pc`, is 32 bits long, as addresses are 32-bit.

Here are the special-purpose registers:

- `flags`: arithmetic/logic FLAGS; (reg encoding: `0x0`)
- `ids`: Interrupt DeStination: the address to jump to upon an interrupt being serviced (also known as the interrupt vector); (reg encoding: `0x1`)
- `ira`: Interrupt Return Address: the address that was jumped from upon an interrupt being serviced; (reg encoding: `0x2`)
- `ie`: IRQ Enable flag: flag indicating whether IRQs are disabled (`0x0`) or enabled (`0x1`); note that this flag starts with a value of `0x0`; (reg encoding: `0x3`)
- `ity`: Interrupt TYpe: flag indicating whether the most recently taken interrupt is an IRQ (`0x0`) or a `swi` (`0x1`) (reg encoding: `0x4`)
- `sty`: Software interrupt TYpe: `swi`'s argument. For `swi rA, #simm`, this is the value `rA + simm` (reg encoding: `0x5`)
- Note: All other encodings for special-purpose registers are reserved.

Here are the bits of `flags`:

- Zero (Z): (`flags` bit 0)
- Carry (C): (`flags` bit 1)
- oVerflow (V): (`flags` bit 2)
- Negative (N): (`flags` bit 3)
- Note: All other bit positions of `flags` are reserved.

II Instruction Set

II.A Instruction Group 0: `pre`, `lpre`, and `atomics`

For `pre`, the following encoding is used, with each character representing one bit:

`0000 iiii iiii iiii`, where

- `i` is a 12-bit constant.

For `lpre`, the following encoding is used, with each character representing one bit:

`0001 0iii iiii iiii iiii iiii iiii iiii`, where

- `i` is a 27-bit constant.

For `cmpxchg` and `xchg`, the following encoding is used, with each character representing one bit:

`0001 100l bbbb aaaa`, where

- `l` encodes whether to lock (1) or not lock (0)
- `b` encodes register `rB`
- `a` encodes register `rA`

This instruction acts as `cmpxchg` if `index` is in effect, using `index_reg` as the expected (or `rC`) value, but otherwise this instruction acts as `xchg`. More on `index` later in this document.

The assembly syntax of these instructions is as follows:

- For `l=0` (without lock):
 - `cmpxchg [rA], rC, rB`
 - `xchg rA, rB`
- For `l=1` (with lock):
 - `cmpxchg.l [rA], rC, rB`
 - `xchg.l [rA], rB`

`cmpxchg/cmpxchg.l` sets the Z flag to 1 upon success and 0 upon failure.

II.A.1 pre and lpre

`pre` and `lpre` are mechanisms by which immediates larger than normal can be used, essentially acting like variable width instructions.

There is no mechanism in the assembly language itself to use `pre` or `lpre` as instructions. Instead, it is expected that the assembler or linker will be the one to insert `pre` or `lpre` as needed if an immediate is too large for a particular instruction.

For non-branch instructions:

- If `pre` is used, the immediate field of the `pre` instruction will form bits [16:5] of the immediate of the next non-`index` instruction. The 17-bit immediate will then be sign-extended to 32 bits. Sign-extension will be performed on the 17-bit immediate even if, had there been no `pre`, the 5-bit immediate would have been zero-extended.
- On the other hand, if `lpre` is used, the immediate field of the `lpre` instruction will form bits [31:5] of the immediate of the next non-`index` instruction.

For branch instructions (group 3):

- if `pre` is used, the immediate field of the `pre` instruction will form bits [20:9] of the immediate of the next non-`index` instruction. The 21-bit immediate will then be sign-extended to 32 bits.
- On the other hand, if `lpre` is used, the immediate field of the `lpre` instruction will form bits [31:9] of the immediate of the next non-`index` instruction.

II.A.2 Handling of pre, lpre, and index

When a `pre` or `lpre` instruction is found, `pre` or `lpre` will be considered to be “in effect”. This condition lasts for one or two instructions after the `pre` or `lpre` instruction, depending on whether or not `index` was in effect.

`index` is an instruction (defined later) that allows a subsequent load or store instruction to perform `base_reg + index_reg` indexing. When an `index` instruction is found, it is considered to be in effect. Like `pre` and `lpre`, `index` is inserted automatically by the assembler.

`index` can be combined with `pre` or `lpre`, and it does not matter whether `index` or the `pre/lpre` instruction came first.

When `pre`, `lpre`, or `index` is in effect, IRQs will not be serviced.

Pseudo code for handling the how `pre`, `lpre`, and `index` are to be handled regarding whether or not they are “in effect” is as follows:

```
function handlePreLpreIndex(input in, output out) {
  when (!out.instr.isLpre()) {
    when (in.instr.isPre()) {
      when (
        in.state.pre.haveIt
        || in.state.lpre.haveIt
      ) {
        // invalid instruction
        out.state.pre.haveIt := False
        out.state.lpre.haveIt := False
        out.state.index.haveIt := False
      }
    }
  }
}
```

```

        out.instr := NOP
        out.state.canServiceIrq := True
    } otherwise {
        out.state.pre.haveIt := True
        out.state.pre.data := in.instr[11:0]
        out.instr := in.instr
        out.state.canServiceIrq := False
    }
} elseif (in.instr.isLpre()) {
    when (
        in.state.pre.haveIt
        || in.state.lpre.haveIt
    ) {
        // invalid instruction
        out.state.pre.haveIt := False
        out.state.lpre.haveIt := False
        out.state.index.haveIt := False
        out.instr := NOP
        out.state.canServiceIrq := True
    } otherwise {
        out.state.lpre.haveIt := True
        out.instr := in.instr
        out.state.canServiceIrq := False
        out.state.lpre.data[26:16] := in.instr[10:0]
    }
} elseif (in.instr.isIndex()) {
    when (
        in.state.index.haveIt
    ) {
        // invalid instruction
        out.state.pre.haveIt := False
        out.state.lpre.haveIt := False
        out.state.index.haveIt := False
        out.instr := NOP
        out.state.canServiceIrq := True
    } otherwise {
        out.state.index.haveIt := True
        out.instr := in.instr
        out.state.canServiceIrq := False
    }
} otherwise {
    out.state.canServiceIrq := !(
        in.state.pre.haveIt
        || in.state.lpre.haveIt
        || in.state.index.haveIt
    )
    out.instr := instr

    // Whenever we see an instruction other than `pre`, `lpre`, or `index`,
    // that means those instructions stop being "in effect".
    out.state.pre.haveIt := False
    out.state.lpre.haveIt := False
    out.state.index.haveIt := False
}
} otherwise { // when (out.instr.isLpre())
    out.state.lpre.data[15:0] := in.instr
}
}

```

II.B Instruction Group 1

The following encoding is used, with each character representing one bit:
001i iiii oooo aaaa, where

- *i* is a 5-bit sign-extended or zero-extended immediate, and is denoted *simm* when sign-extended or *imm* when zero-extended. Also, *simm* or *imm* can be expanded with *pre* or *lpre*.
- *a* encodes register *rA*
- *o* is the opcode

Here is a list of instructions from this encoding group.

- Opcode 0x0: `add rA, #simm`
- Opcode 0x1: `add rA, pc, #simm`
 - Effect: `rA <= pc + simm + 2;`
- Opcode 0x2: `add rA, sp, #simm`
- Opcode 0x3: `add rA, fp, #simm`
- Opcode 0x4: `cmp rA, #simm`
 - Effect: Compare *rA* to *simm*.
 - Affectable flags: Z, C, V, N
- Opcode 0x5: `cpy rA, #simm`
 - Effect: Copy an immediate value into *rA*
- Opcode 0x6: `lsl rA, #imm`
 - Effect: Logical shift left
- Opcode 0x7: `lsr rA, #imm`
 - Effect: Logical shift right
- Opcode 0x8: `asr rA, #imm`
 - Effect: Arithmetic shift right
- Opcode 0x9: `and rA, #simm`
 - Effect: Bitwise AND
- Opcode 0xa: `orr rA, #simm`
 - Effect: Bitwise OR
- Opcode 0xb: `xor rA, #simm`
 - Effect: Bitwise XOR
- Opcode 0xc: `ze rA, #imm`
 - Effect: Set *rA*[31:imm] to zero.
- Opcode 0xd: `se rA, #imm`
 - Effect: Set each bit of *rA*[31:imm] to the bit *rA*[imm].
- Opcode 0xe: `swi rA, #simm`
 - Effect: Call software interrupt number *rA + simm*.
- Opcode 0xf: `swi #imm`
 - Effect: Call software interrupt number *imm*.

II.C Instruction Group 2

The following encoding is used, with each character representing one bit:

010f oooo bbbb aaaa, where

- *o* is the opcode
- *b* encodes register *rB*
- *a* encodes register *rA*
- *f* is encoded as 0 if this instruction cannot affect flags and encoded 1 if this instruction is permitted to affect flags. Note that `cmp` is permitted to affect flags regardless of this bit.

Here is a list of instructions from this encoding group.

- Opcode 0x0: `add rA, rB`
 - Mnemonic for when flags not affected: `add`
 - Mnemonic for when flags affected: `add.f`
 - Affectable flags: Z, C, V, N
- Opcode 0x1: `sub rA, rB`
 - Mnemonic for when flags not affected: `sub`
 - Mnemonic for when flags affected: `sub.f`

- Affectable flags: Z, C, V, N
- Opcode 0x2: `add rA, sp, rB`
 - Mnemonic for when flags not affected: `add`
 - Mnemonic for when flags affected: `add.f`
 - Affectable flags: Z, C, V, N
- Opcode 0x3: `add rA, fp, rB`
 - Mnemonic for when flags not affected: `add`
 - Mnemonic for when flags affected: `add.f`
 - Affectable flags: Z, C, V, N
- Opcode 0x4: `cmp rA, rB`
 - Effect: Compare `rA` to `rB`. `cmp` is *always* able to affect flags, independent of the encoded `f` bit of the instruction.
 - Affectable flags: Z, C, V, N
- Opcode 0x5: `cpy rA, rB`
 - Mnemonic for when flags not affected: `cpy`
 - Mnemonic for when flags affected: `cpy.f`
 - Effect: Copy `rB` into `rA`
 - Affectable flags: Z, N
- Opcode 0x6: `lsl rA, rB`
 - Mnemonic for when flags not affected: `lsl`
 - Mnemonic for when flags affected: `lsl.f`
 - Effect: Logical shift left
 - Affectable flags: Z, N
- Opcode 0x7: `lsr rA, rB`
 - Mnemonic for when flags not affected: `lsr`
 - Mnemonic for when flags affected: `lsr.f`
 - Effect: Logical shift right
 - Affectable flags: Z, N
- Opcode 0x8: `asr rA, rB`
 - Mnemonic for when flags not affected: `asr`
 - Mnemonic for when flags affected: `asr.f`
 - Effect: Arithmetic shift right
 - Affectable flags: Z, N
- Opcode 0x9: `and rA, rB`
 - Mnemonic for when flags not affected: `and`
 - Mnemonic for when flags affected: `and.f`
 - Effect: Bitwise AND
 - Affectable flags: Z, N
- Opcode 0xa: `orr rA, rB`
 - Mnemonic for when flags not affected: `orr`
 - Mnemonic for when flags affected: `orr.f`
 - Effect: Bitwise OR
 - Affectable flags: Z, N
- Opcode 0xb: `xor rA, rB`
 - Mnemonic for when flags not affected: `xor`
 - Mnemonic for when flags affected: `xor.f`
 - Effect: Bitwise XOR
 - Affectable flags: Z, N
- Opcode 0xc: `adc rA, rB`
 - Mnemonic for when flags not affected: `adc`
 - Mnemonic for when flags affected: `adc.f`
 - Effect: Add with Carry, using the formula `rA + rB + flags.C` to compute the value that will be written into `rA`.

- Affectable flags: Z, C, V, N
- Opcode 0xd: sbc rA, rB
 - Mnemonic for when flags not affected: sbc
 - Mnemonic for when flags affected: sbc.f
 - Effect: Subtract with Borrow, using the formula $rA + (\sim rB) + \text{flags.C}$ to compute the value that will be written into rA.
 - Affectable flags: Z, C, V, N
- Opcode 0xe: cmpbc rA, rB
 - Effect: Compare rA to rB, but with carry-in and a different effect for setting the flags.Z. cmpbc is *always* able to affect flags, independent of the encoded f bit of the instruction.
 - Note: this instruction acts much like sbc rA, rB, but without storing the subtraction's result into rA. However, this instruction sets the Z flag to $\text{prev}(\text{flags.Z}) \text{ AND } ((rA + (\sim rB) + \text{flags.C}) == 0)$
 - Affectable flags: Z, C, V, N

II.D Instruction Group 3: Relative Branches

The following encoding is used, with each character representing one bit:

011i iiii iiii oooo, where

- i is a 9-bit sign-extended immediate, which can be expanded by pre or lpre, and is denoted *simm*
- o is the opcode

Here is a list of instructions from this encoding group.

- Opcode 0x0: bl *simm*
 - Name: Branch and Link
 - Description: Relative call
 - Effect: $lr \leq pc + 2$; $pc \leq pc + \text{simm} + 2$;
- Opcode 0x1: bra *simm*
 - Name: BRanch Always
 - Description: Unconditional relative branch
 - Effect: $pc \leq pc + \text{simm} + 2$;
- Opcode 0x2: beq *simm*
 - Name: Branch if EQual
 - Effect: if (flags.Z) $pc \leq pc + \text{simm} + 2$;
- Opcode 0x3: bne *simm*
 - Name: Branch if Not Equal
 - Effect: if (!flags.Z) $pc \leq pc + \text{simm} + 2$;
- Opcode 0x4: bmi *simm*
 - Name: Branch if MInus
 - Effect: if (flags.N) $pc \leq pc + \text{simm} + 2$;
- Opcode 0x5: bpl *simm*
 - Name: Branch if PPlus
 - Effect: if (!flags.N) $pc \leq pc + \text{simm} + 2$;
- Opcode 0x6: bvs *simm*
 - Name: Branch if oVerflow Set
 - Effect: if (flags.V) $pc \leq pc + \text{simm} + 2$;
- Opcode 0x7: bvc *simm*
 - Name: Branch if oVerflow Clear
 - Effect: if (!flags.V) $pc \leq pc + \text{simm} + 2$;
- Opcode 0x8: bgeu *simm*
 - Name: Branch if Greater than or Equal Unsigned
 - Effect: if (flags.C) $pc \leq pc + \text{simm} + 2$;
- Opcode 0x9: bltu *simm*
 - Name: Branch if Less Than Unsigned

- Effect: if (!flags.C) pc <= pc + simm + 2;
- Opcode 0xa: bgtu simm
 - Name: Branch if Greater Than Unsigned
 - Effect: if (flags.C AND !flags.Z) pc <= pc + simm + 2;
- Opcode 0xb: bleu simm
 - Name: Branch if Less than or Equal Unsigned
 - Effect: if (!flags.C OR flags.Z) pc <= pc + simm + 2;
- Opcode 0xc: bges simm
 - Name: Branch if Greater than or Equal Signed
 - Effect: if (!(flags.N XOR flags.V)) pc <= pc + simm + 2;
- Opcode 0xd: blts simm
 - Name: Branch if Less Than Signed
 - Effect: if (flags.N XOR flags.V) pc <= pc + simm + 2;
- Opcode 0xe: bgts simm
 - Name: Branch if Greater Than Signed
 - Effect:


```
if (!(flags.N XOR flags.V) AND !flags.Z)
    pc <= pc + simm + 2;
```
- Opcode 0xf: bles simm
 - Name: Branch if Less than or Equal Signed
 - Effect:


```
if ((flags.N XOR flags.V) OR flags.Z)
    pc <= pc + simm + 2;
```

II.E Instruction Group 4

The following encoding is used, with each character representing one bit:

1000 0000 bbbb aaaa, where

- o is the opcode
- b encodes register rB or register sC, where sC is one of the special-purpose registers
- a encodes register rA or register sA, where sA is one of the special-purpose registers

Here is a list of instructions from this encoding group.

- Opcode 0x0: jl rA
 - Effect: lr <= pc + 2; pc <= rA;
- Opcode 0x1: jmp rA
 - Effect: pc <= rA;
- Opcode 0x2: jmp ira
 - Effect: pc <= ira;
- Opcode 0x3: reti
 - Effect: enables IRQs (by copying 0x1 into ie) and performs pc <= ira;
- Opcode 0x4: ei
 - Effect: copy 1 into ie.
- Opcode 0x5: di
 - Effect: copy 0 into ie.
- Opcode 0x6: push rA, rB
 - Effect: pushes rA onto the stack, using rB as the stack pointer, post-decrementing rB.
 - This instruction does nothing when rA is the same register as rB.
 - Note: As a pseudo instruction, omitting “, rB” will automatically select sp as the particular stack pointer.
- Opcode 0x7: push sA, rB
 - Effect: pushes sA onto the stack, using rB as the stack pointer, post-decrementing rB.
 - Note that sA is considered to be 32-bit for the purpose of the store to memory and decrementing rB, even if sA is flags or ie.

- Note: As a pseudo instruction, omitting “, rB” will automatically select **sp** as the particular stack pointer.
- Opcode 0x8: **pop rA, rB**
 - Effect: pops **rA** off the stack, using **rB** as the stack pointer, pre-incrementing **rB**.
 - This instruction does nothing when **rA** is the same register as **rB**.
 - Note: As a pseudo instruction, omitting “, rB” will automatically select **sp** as the particular stack pointer.
- Opcode 0x9: **pop sA, rB**
 - Effect: pops **sA** off the stack, using **rB** as the stack pointer, pre-incrementing **rB**.
 - Note that **sA** is considered to be 32-bit for the purpose of the load from memory and incrementing **rB**, even if **sA** is **flags** or **ie**.
 - Note: As a pseudo instruction, omitting “, rB” will automatically select **sp** as the particular stack pointer.
- Opcode 0xa: **pop pc, rB**
 - Effect: pops **pc** off the stack, using **rB** as the stack pointer, pre-incrementing **rB**.
 - Note: As a pseudo instruction, omitting “, rB” will automatically select **sp** as the particular stack pointer.
- Opcode 0xb: **mul rA, rB**
 - Effect: $rA \leq rA * rB$;
- Opcode 0xc: **udiv rA, rB**
 - Effect: $rA \leq u32(rA) / u32(rB)$;
 - Note: This instruction executes more quickly if the values of **rA** and **rB** that are input to this instruction were the same for the previous 32-bit divide or modulo instruction.
- Opcode 0xd: **sdiv rA, rB**
 - Effect: $rA \leq s32(rA) / s32(rB)$;
 - Note: This instruction executes more quickly if the values of **rA** and **rB** that are input to this instruction were the same for the previous 32-bit divide or modulo instruction.
- Opcode 0xe: **umod rA, rB**
 - Effect: $rA \leq u32(rA) \% u32(rB)$;
 - Note: This instruction executes more quickly if the values of **rA** and **rB** that are input to this instruction were the same for the previous 32-bit divide or modulo instruction.
- Opcode 0xf: **smod rA, rB**
 - Effect: $rA \leq s32(rA) \% s32(rB)$;
 - Note: This instruction executes more quickly if the values of **rA** and **rB** that are input to this instruction were the same for the previous 32-bit divide or modulo instruction.
- Opcode 0x10: **lumul rA, rB**
 - Effect: This instruction multiplies **rA** by **rB**, performing an unsigned 32-bit by 32-bit -> 64-bit multiply, storing result in **concat{r0, r1}**.
- Opcode 0x11: **lsmul rA, rB**
 - Effect: This instruction multiplies **rA** by **rB**, performing a signed 32-bit by 32-bit -> 64-bit multiply, storing result in **concat{r0, r1}**.
- Opcode 0x12: **udiv64 rA, rB**
 - Effect: performs a 64-bit by 64-bit unsigned division of **concat{rA, r{A + 1}}** by **concat{rB, r{B + 1}}**, storing 64-bit result in **concat{rA, r{A + 1}}**.
 - Note: This instruction operates as if **rA** were encoded with (**A[0x0] == 0b0**) and **rB** were encoded with (**B[0x0] == 0b0**),
 - Note: This instruction executes more quickly if **rB**'s value is **0x00000000**, i.e. if the operation is actually a 64-bit by 32-bit -> 64-bit unsigned divide.
 - Note: This instruction executes more quickly if the values of **concat{rA, r{A + 1}}** and **concat{rB, r{B + 1}}** that are input to this instruction were the same for the previous 64-bit divide or modulo instruction.
- Opcode 0x13: **sdiv64 rA, rB**

- Effect: performs a 64-bit by 64-bit signed division of `concat{rA, r{A + 1}}` by `concat{rB, r{B + 1}}`, storing 64-bit result in `concat{rA, r{A + 1}}`.
- Note: This instruction operates as if rA were encoded with (`A[0x0] == 0b0`) and rB were encoded with (`B[0x0] == 0b0`),
- Note: This instruction executes more quickly if rB is equal to bits [63:32] of `s64(r{B+1})`.
- Note: This instruction executes more quickly if the values of `concat{rA, r{A + 1}}` and `concat{rB, r{B + 1}}` that are input to this instruction were the same for the previous 64-bit divide or modulo instruction.
- Opcode 0x14: `umod64 rA, rB`
 - Effect: performs a 64-bit by 64-bit unsigned modulo of `concat{rA, r{A + 1}}` by `concat{rB, r{B + 1}}`, storing 64-bit result in `concat{rA, r{A + 1}}`.
 - Note: This instruction operates as if rA were encoded with (`A[0x0] == 0b0`) and rB were encoded with (`B[0x0] == 0b0`),
 - Note: This instruction executes more quickly if rB's value is `0x00000000`, i.e. if the operation is actually a 64-bit by 32-bit -> 64-bit unsigned modulo.
 - Note: This instruction executes more quickly if the values of `concat{rA, r{A + 1}}` and `concat{rB, r{B + 1}}` that are input to this instruction were the same for the previous 64-bit divide or modulo instruction.
- Opcode 0x15: `smod64 rA, rB`
 - Effect: performs a 64-bit by 64-bit signed modulo of `concat{rA, r{A + 1}}` by `concat{rB, r{B + 1}}`, storing 64-bit result in `concat{rA, r{A + 1}}`.
 - Note: This instruction operates as if rA were encoded with (`A[0x0] == 0b0`) and rB were encoded with (`B[0x0] == 0b0`),
 - Note: This instruction executes more quickly if rB is equal to bits [63:32] of `s64(r{B+1})`.
 - Note: This instruction executes more quickly if the values of `concat{rA, r{A + 1}}` and `concat{rB, r{B + 1}}` that are input to this instruction were the same for the previous 64-bit divide or modulo instruction.
- Opcode 0x16: `ldub rA, [rB]`
 - Effect: Load an 8-bit value from memory at address computed as `rB + index_reg`, zero-extend the 8-bit value to 32 bits, then put the zero-extended 32-bit value into rA.
 - The `index_reg` value is guaranteed to be zero unless an `index` is in effect.
 - Shorthand for having the assembler insert an `index rC` instruction before this one: `ldub rA, [rB, rC]`
- Opcode 0x17: `ldsb rA, [rB]`
 - Effect: Load an 8-bit value from memory at address computed as `rB + index_reg`, sign-extend the 8-bit value to 32 bits, then put the sign-extended 32-bit value into rA.
 - The `index_reg` value is guaranteed to be zero unless an `index` is in effect.
 - Shorthand for having the assembler insert an `index rC` instruction before this one: `ldsb rA, [rB, rC]`
- Opcode 0x18: `lduh rA, [rB]`
 - Effect: Load a 16-bit value from memory at address computed as `rB + index_reg`, zero-extend the 16-bit value to 32 bits, then put the zero-extended 32-bit value into rA.
 - The `index_reg` value is guaranteed to be zero unless an `index` is in effect.
 - Shorthand for having the assembler insert an `index rC` instruction before this one: `lduh rA, [rB, rC]`
- Opcode 0x19: `ldsh rA, [rB]`
 - Effect: Load a 16-bit value from memory at address computed as `rB + index_reg`, sign-extend the 16-bit value to 32 bits, then put the zero-extended 32-bit value into rA.
 - The `index_reg` value is guaranteed to be zero unless an `index` is in effect.
 - Shorthand for having the assembler insert an `index rC` instruction before this one: `ldsh rA, [rB, rC]`
- Opcode 0x1a: `stb rA, [rB]`
 - Effect: Store `rA[7:0]` to memory at the address computed as `rB + index_reg`.

- The `index_reg` value is guaranteed to be zero unless an `index` is in effect.
- Shorthand for having the assembler insert an `index rC` instruction before this one: `stb rA, [rB, rC]`
- Opcode `0x1b`: `sth rA, [rB]`
 - Effect: Store `rA[15:0]` to memory at the address computed as `rB + index_reg`.
 - The `index_reg` value is guaranteed to be zero unless an `index` is in effect.
 - Shorthand for having the assembler insert an `index rC` instruction before this one: `sth rA, [rB, rC]`
- Opcode `0x1c`: `cpy rA, sB`
 - Effect: `rA <= sB`;
- Opcode `0x1d`: `cpy sA, rB`
 - Effect: `sA <= rB`;
- Opcode `0x1e`: `cpy sA, sB`
 - Effect: `sA <= sB`;
- Opcode `0x1f`: `index rA`
 - Effect: Performs `index_reg <= rA`; and stores that `index` is in effect.
 - Note: If `index` is in effect and the current instruction is `index`, the current instruction will be treated as a NOP, and `index` will stop being in effect.
 - Note: `pre` and `index` can be combined with one another (though this is only useful for `ldr` and `str`).
 - Note: A non-`pre` instruction following `index` will store that that `index` is not in effect any more. (It will also store that `pre` is not in effect any more).
 - Note: If `index` is in effect, the current instruction cannot be interrupted by an IRQ.
 - Note: Any time `index` stops being in effect, `pre` will stop being in effect as well.

II.F Instruction Group 5: Immediate Indexed Load

The following encoding is used, with each character representing one bit:

`101i iiii bbbb aaaa`, where

- `i` is a 5-bit sign-extended immediate, which can be expanded by `pre` or `lpre`, and is denoted `sim`
- `b` encodes register `rB`
- `a` encodes register `rA`

The one instruction from this encoding group is `ldr rA, [rB, #sim]`. This is a 32-bit load into `rA`, where the effective address to load from is computed as `rB + <index_reg> + sim`, using the sign-extended form of `sim`.

The `<index_reg>` value is guaranteed to be zero unless an `index` is in effect.

The following pseudo instructions exist within the assembler for this encoding group:

- `ldr rA, [rB]`
 - This pseudo instruction assembles to the following:
 - `ldr rA, [rB, #0x0]`
- `ldr rA, [rB, rC]`
 - This pseudo instruction assembles to the following:
 - `index rC`
 - `ldr rA, [rB, #0x0]`
- `ldr rA, [rB, rC, #sim]`
 - This pseudo instruction assembles to the following:
 - `index rC`
 - `ldr rA, [rB, #sim]`

II.G Instruction Group 6: Immediate Indexed Store

The following encoding is used, with each character representing one bit:

`110i iiii bbbb aaaa`, where

- `i` is a 5-bit sign-extended immediate, which can be expanded by `pre` or `lpre`, and is denoted `sim`

- **b** encodes register **rB**
- **a** encodes register **rA**

The one instruction from this encoding group is **str rA, [rB, #simm]**. This is a 32-bit store of **rA**, where the effective address to store to is computed as **<index_reg> + rB + simm**, using the sign-extended form of **simm**.

The **<index_reg>** value is guaranteed to be zero unless an **index** is in effect.

The following pseudo instructions exist within the assembler for this encoding group:

- **str rA, [rB]**
 - This pseudo instruction assembles to the following:
 - **str rA, [rB, #0x0]**
- **str rA, [rB, rC]**
 - This pseudo instruction assembles to the following:
 - **index rC**
 - **str rA, [rB, #0x0]**
- **str rA, [rB, rC, #simm]**
 - This pseudo instruction assembles to the following:
 - **index rC**
 - **str rA, [rB, #simm]**

II.H Instruction Group 7, Subgroup **0b00**:

Extra 8-bit and 16-bit Ops} The following encoding is used, with each character representing one bit:

1110 0w00 bbbb aaaa, where

- **w** is the operation width
 - When **0b0**: 8-bit operation
 - When **0b1**: 16-bit operation
- **o** is the opcode
- **b** encodes register **rB**
- **a** encodes register **rA**

Here is a list of instructions from this encoding group.

- Opcode **0x0**:
 - **w** value **0b0**: **cmpb rA, rB**
 - Effect: Compare **rA[7:0]** to **rB[7:0]**
 - Affectable flags: **Z, C, V, N**
 - **w** value **0b1**: **cmph rA, rB**
 - Effect: Compare **rA[15:0]** to **rB[15:0]**
 - Affectable flags: **Z, C, V, N**
- Opcode **0x1**:
 - **w** value **0b0**: **lsrb rA, rB**
 - Effect: Logical shift right **rA[7:0]** by **rB**
 - **w** value **0b1**: **lsrh rA, rB**
 - Effect: Logical shift right **rA[15:0]** by **rB**
- Opcode **0x2**:
 - **w** value **0b0**: **asrb rA, rB**
 - Effect: Arithmetic shift right **rA[7:0]** by **rB**
 - **w** value **0b1**: **asrh rA, rB**
 - Effect: Arithmetic shift right **rA[15:0]** by **rB**

II.I Instruction Group 7, Subgroup **0b010**: Extra load/store instructions

The following encoding is used, with each character representing one bit:

1110 1000 bbbb aaaa, where

- **o** is the opcode

- **b** encodes register **rB** or **sB**
- **a** encodes register **sA**

Here is a list of instructions from this encoding group.

- Opcode **0x0**: **ldr sA, [rB]**
 - Effect: 32-bit load of **sA** from memory at address held in **rB**.
- Opcode **0x1**: **ldr sA, [sB]**
 - Effect: 32-bit load of **sA** from memory at address held in **sB**.
- Opcode **0x2**: **str sA, [rB]**
 - Effect: 32-bit store of **sA** to memory at address held in **rB**.
- Opcode **0x3**: **str sA, [sB]**
 - Effect: 32-bit store of **sA** to memory at address held in **sB**.

II.J Instruction Group 7, Subgroup **0b0110: Icache Flush Instruction**

The following encoding is used, with each character representing one bit:

1110 1100 0000 0000

The one instruction from this encoding group is **icflush**, which invalidates the instruction cache entirely.